

A large grid of circular nodes, each containing a small triangle icon, is shown. The nodes are colored in a gradient from blue at the bottom left to red at the top right, suggesting a cost or distance field. A yellow rectangular box in the upper right corner contains the text "Games AI", "Lecture 2.1", and "Pathfinding".

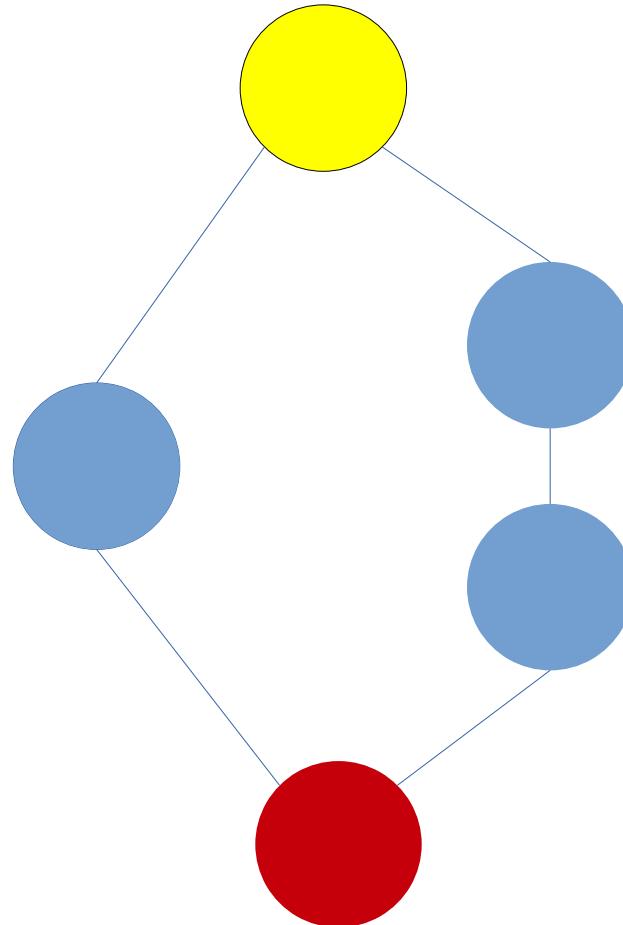
Games AI

Lecture 2.1

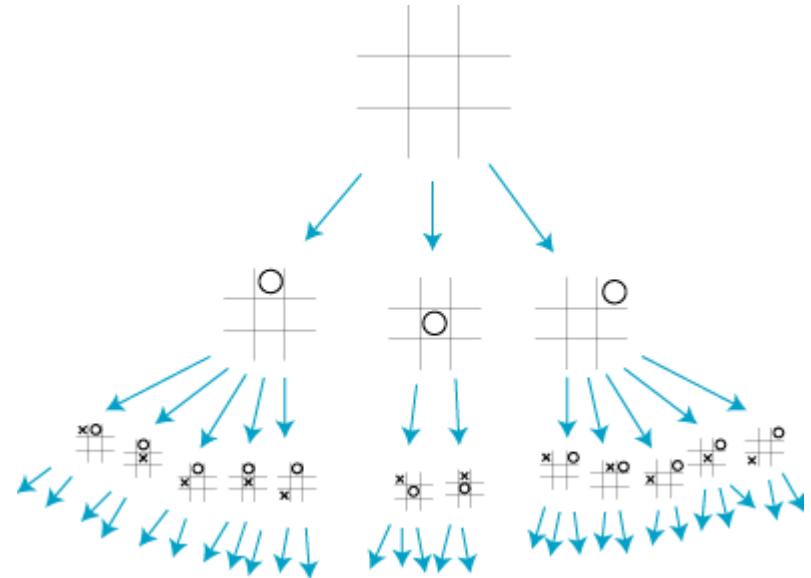
Pathfinding

- Pathfinding
 - Breadth-first search
 - A* search
 - Heuristics
 - Navigation graphs
 - A* variants

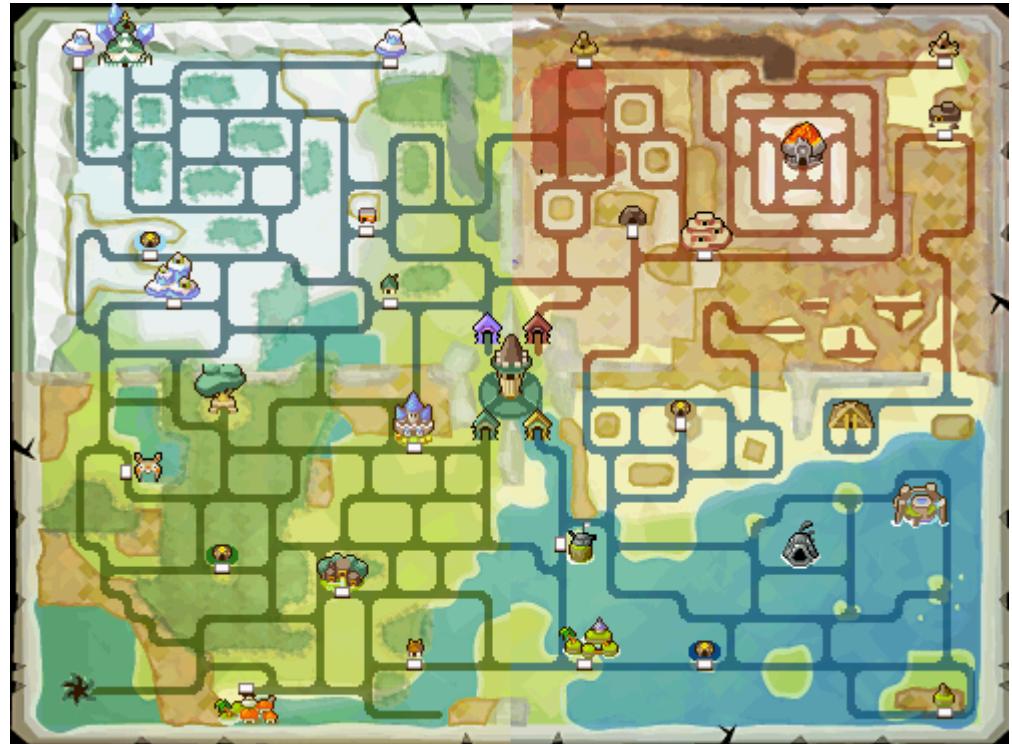
- Graph Search
 - What is a graph?



- Graph Search
 - What is a graph?

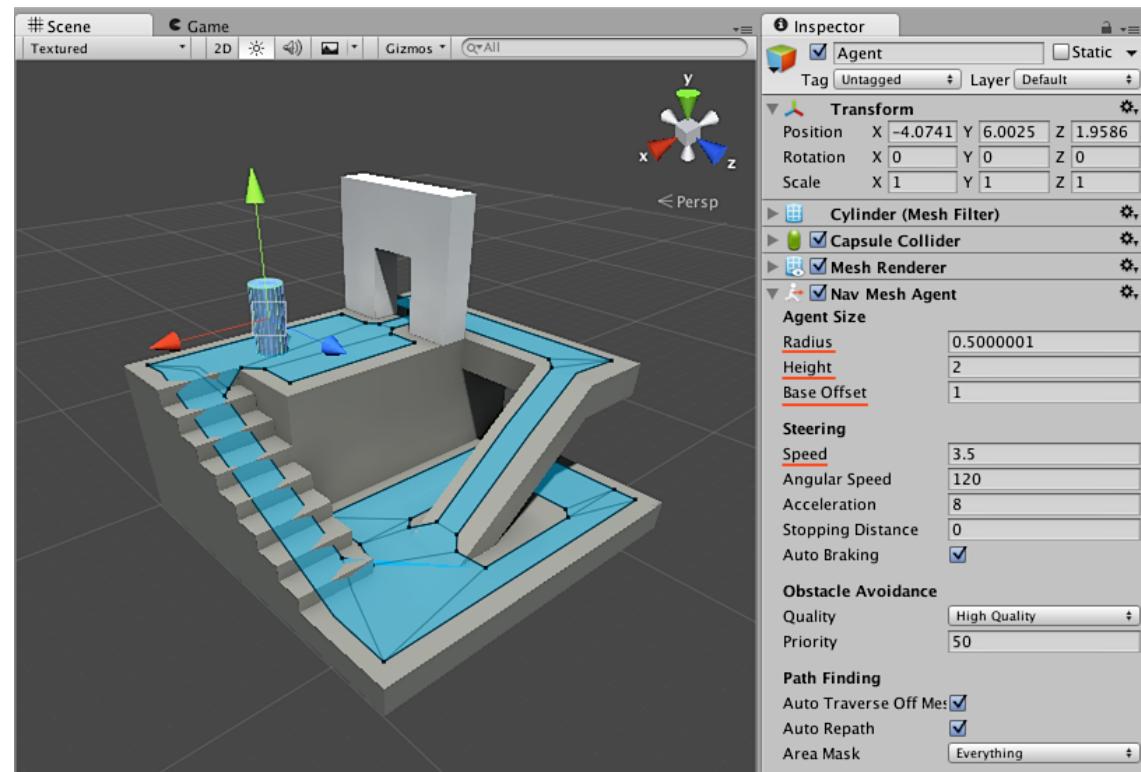


- Graph Search
 - What is a graph?

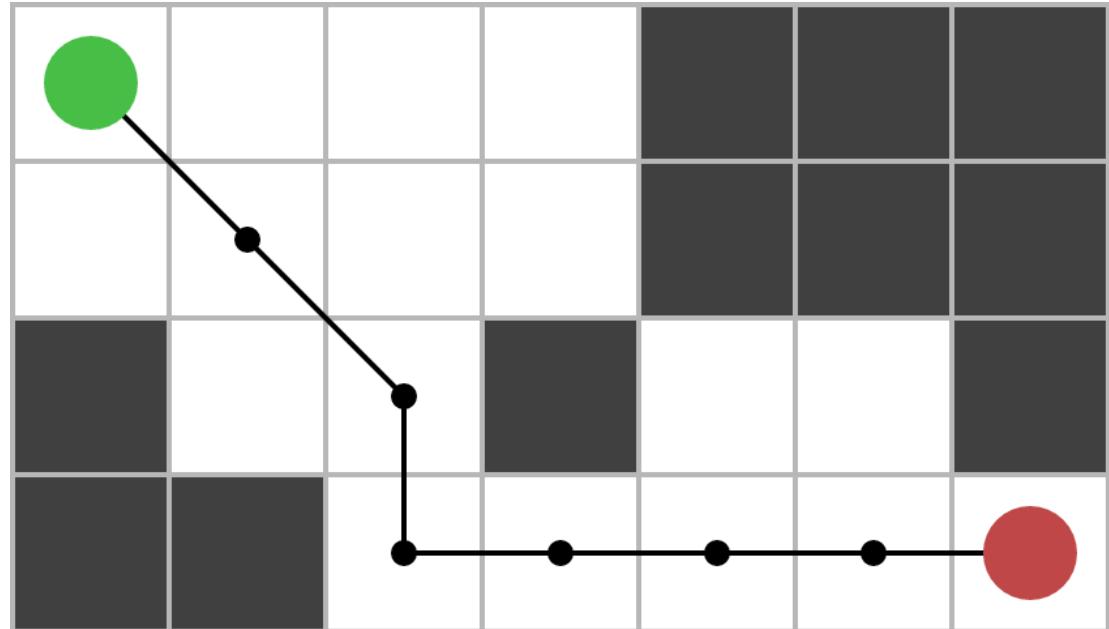


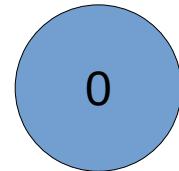
Tree search

- Graph Search
 - What is a graph?



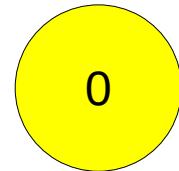
- Graph Search
 - What is a graph?





Queue
0

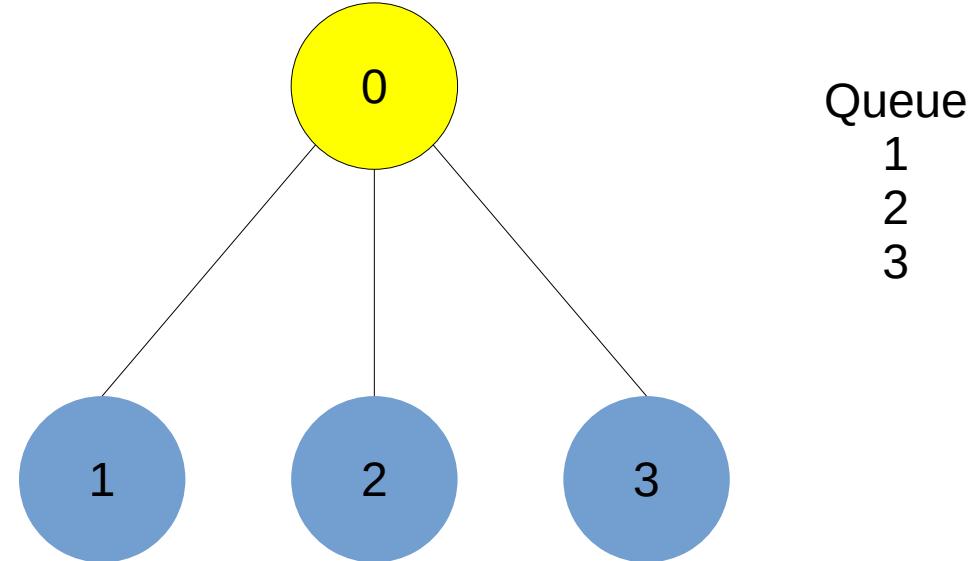
- Breadth-first search
- Set up:
 1. Create queue
 2. Add start to queue
 3. Mark as discovered



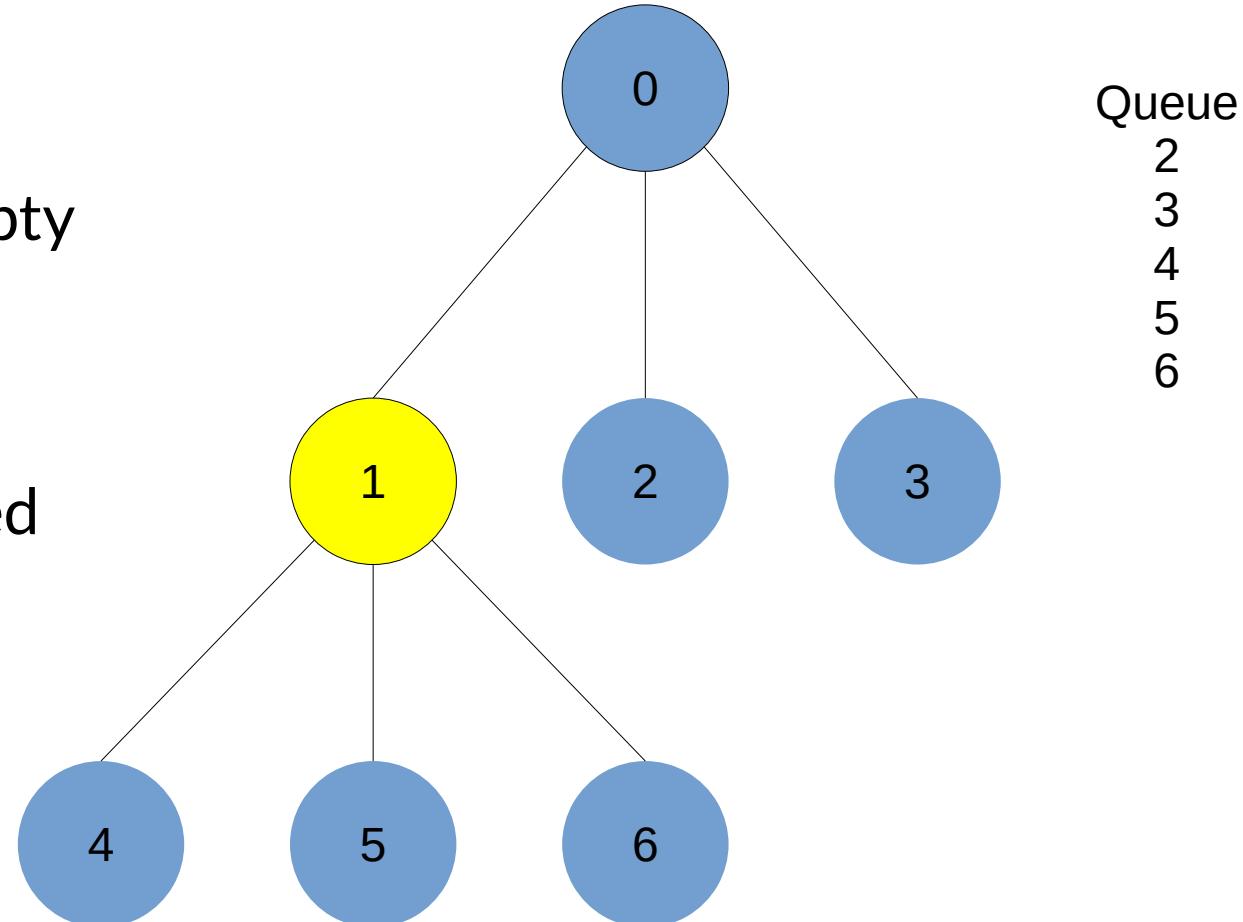
Queue

- While queue is not empty
 - 1. Select top of queue
 - 2. Check if goal

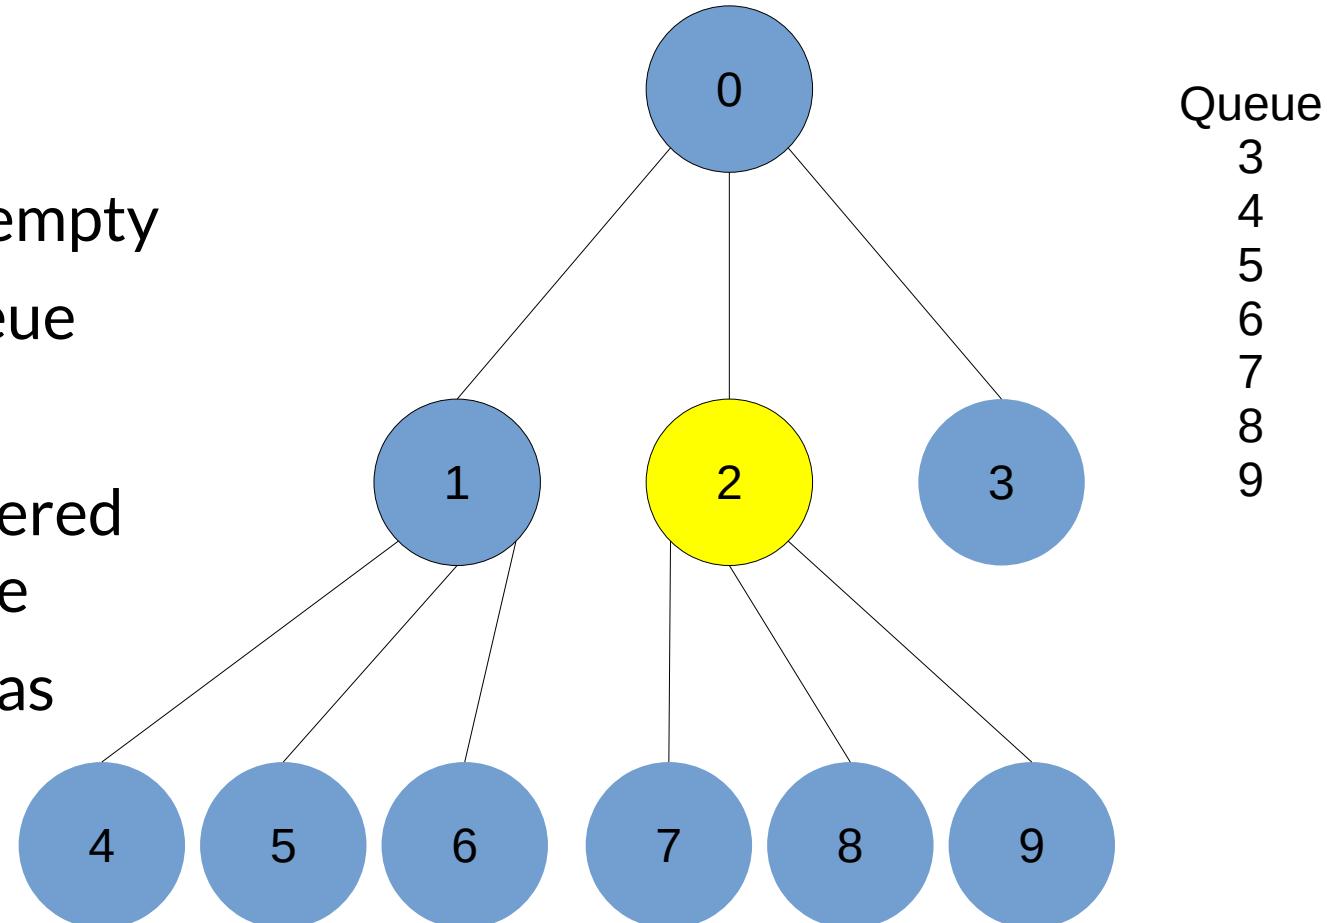
- While queue is not empty
 1. Select top of queue
 2. Check if goal
 3. Add all undiscovered children to queue
 4. Mark each child as discovered



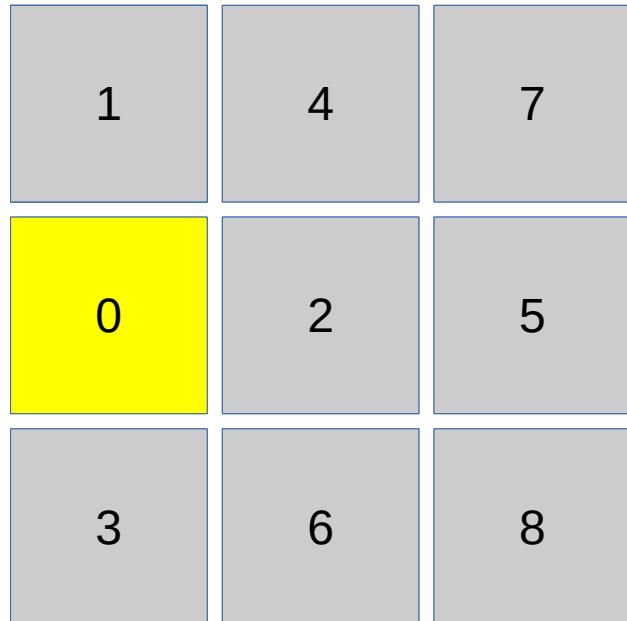
- While queue is not empty
 - 1. Select top of queue
 - 2. Check if goal
 - 3. Add all undiscovered children to queue
 - 4. Mark each child as discovered



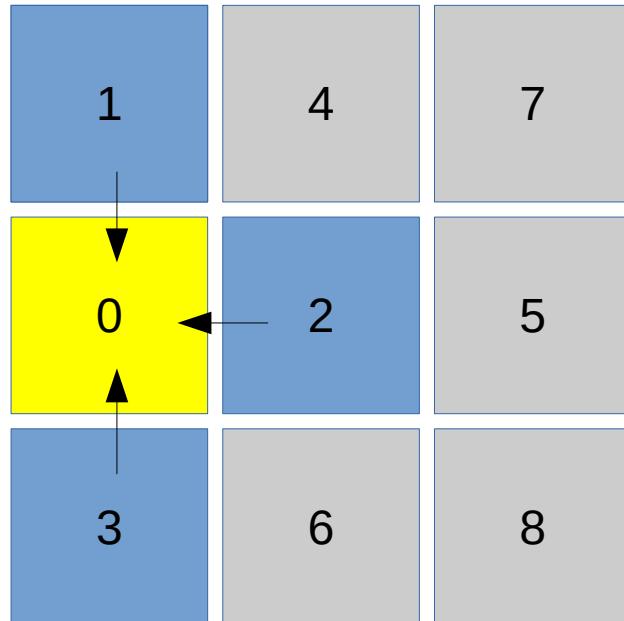
- While queue is not empty
 1. Select top of queue
 2. Check if goal
 3. Add all undiscovered children to queue
 4. Mark each child as discovered



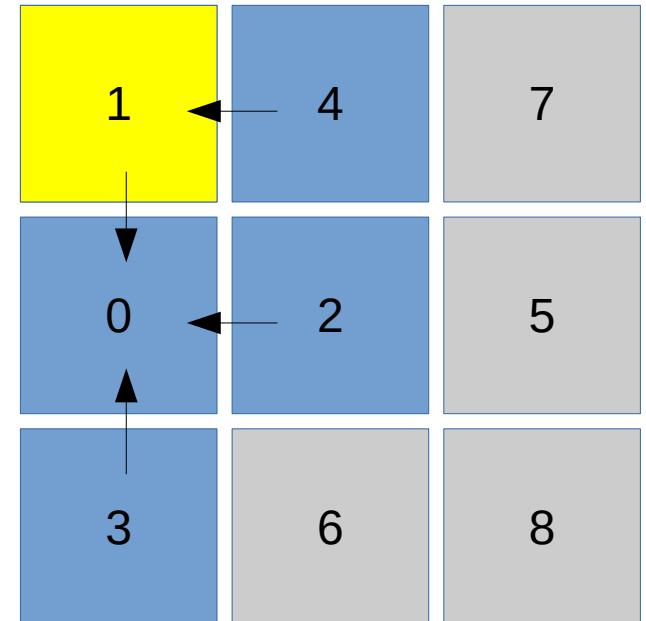
BFS



Queue:

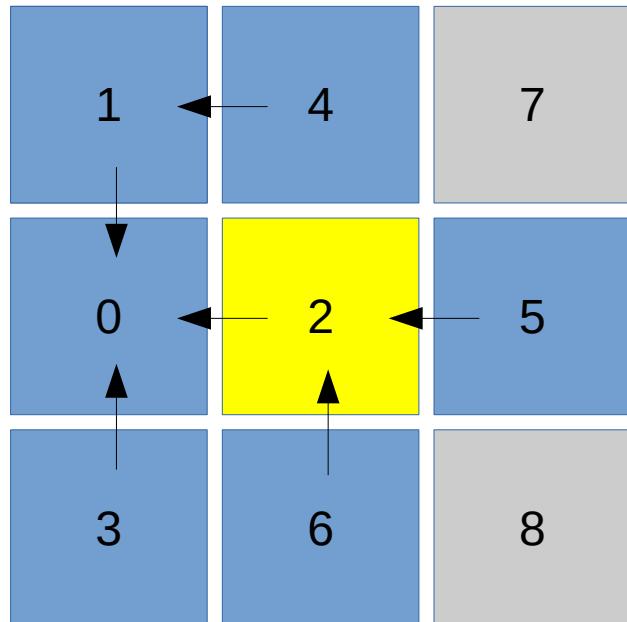


Queue: 1,2,3

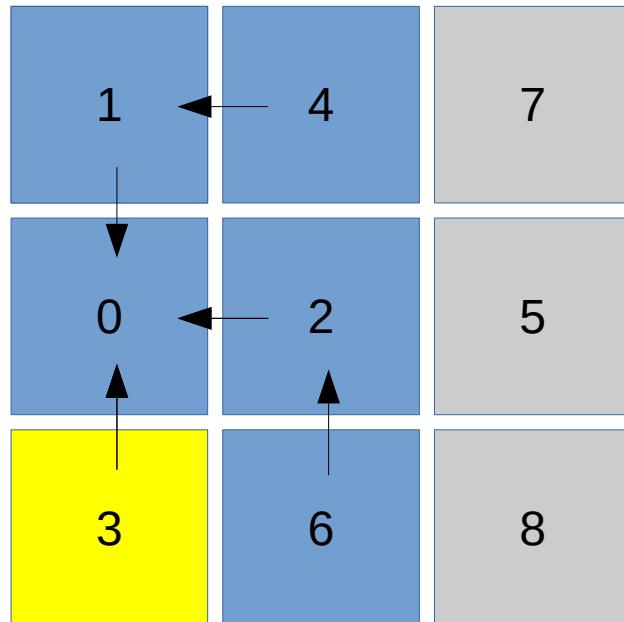


Queue: 2,3,4

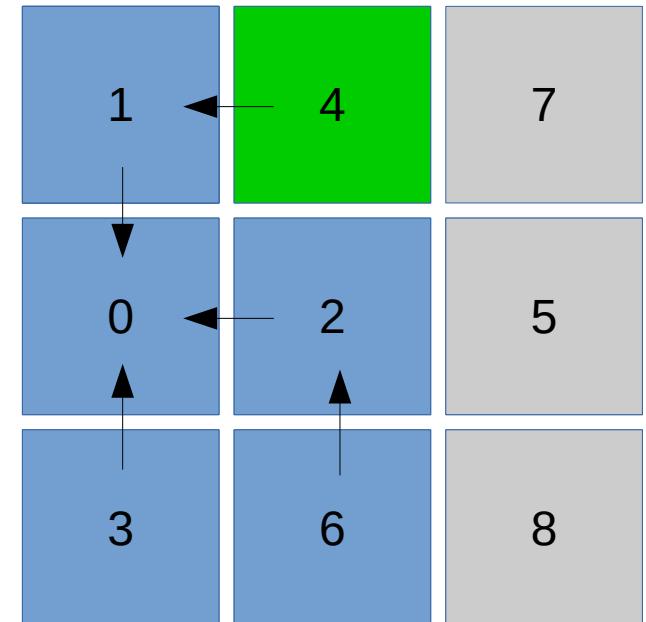
BFS



Queue: 3,4,5,6



Queue: 4,5,6



Queue: 5,6,7

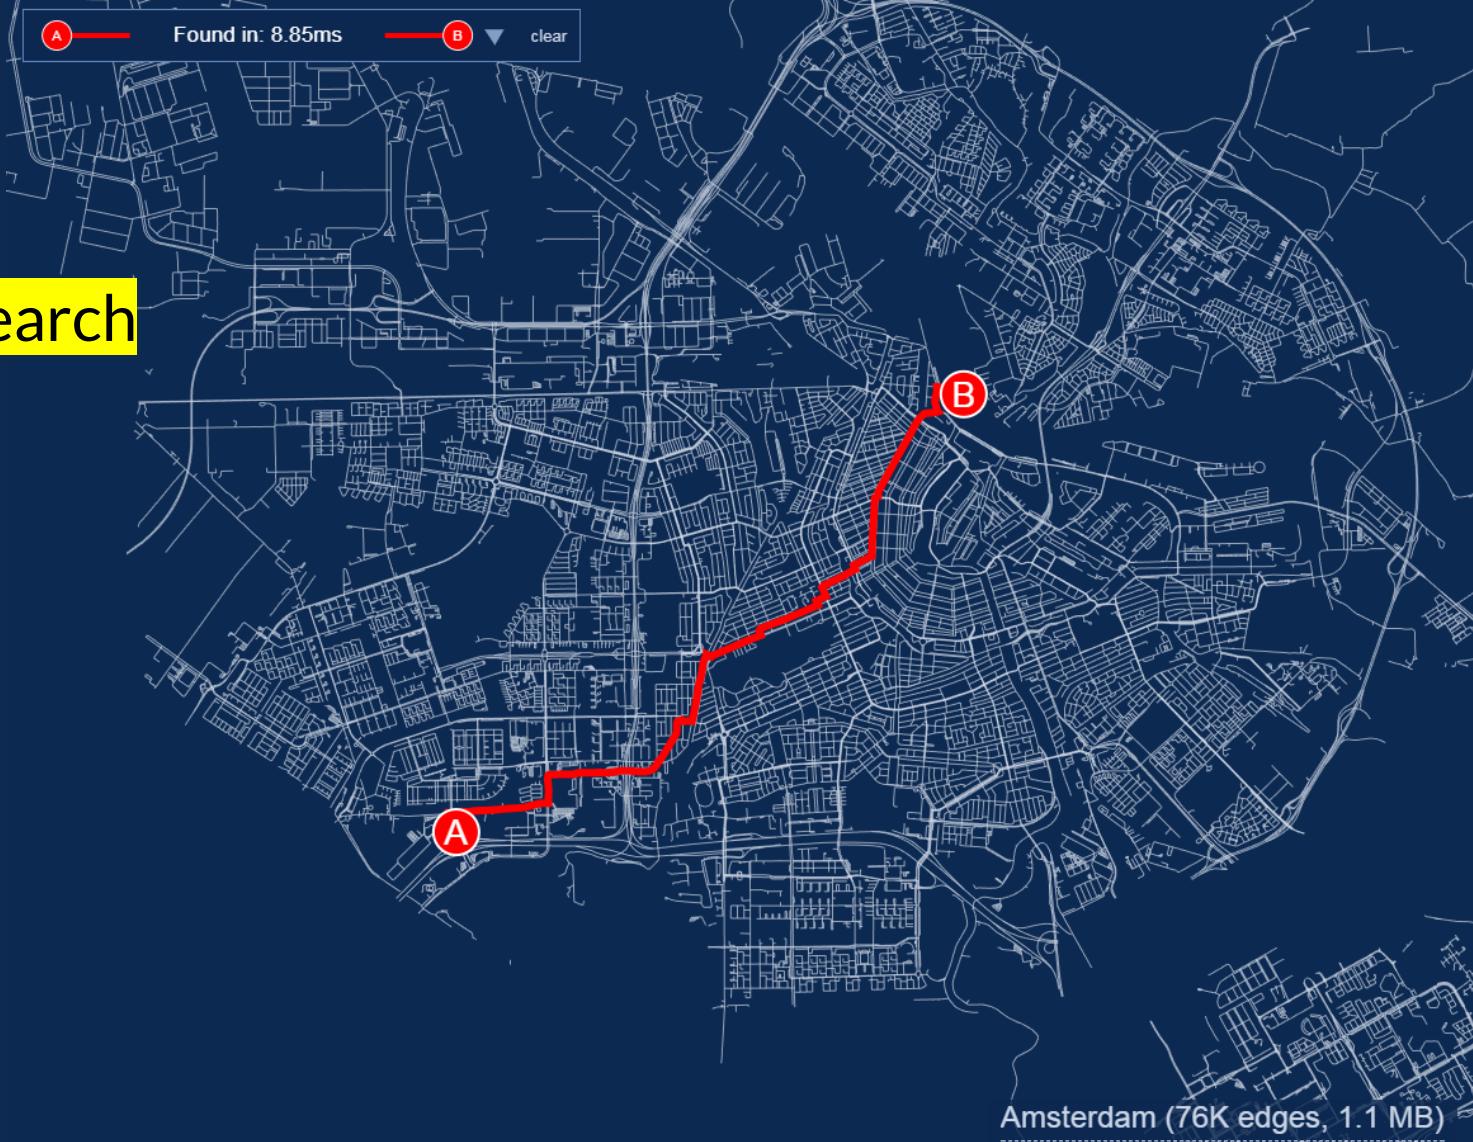
BFS

```
1. procedure BFS(G: Graph, v: Vertex of G) is
2.   create a queue Q
3.   enqueue v onto Q
4.   mark v
5.   while Q is not empty do
6.     w ← Q.dequeue()
7.     if w is what we are looking for then
8.       return w
9.     for all edges e in G.adjacentEdges(w) do
10.      x ← G.adjacentVertex(w, e)
11.      if x is not marked then
12.        mark x
13.        enqueue x onto Q
14.  return nulls.
```

```
breadthFirstSearch(initialState: T): A[]
{
    this.openList.push(new TreeNode<T, A>(initialState, null, null));
    while (this.openList.length > 0)
    {
        let current = this.openList.shift();
        if (this.goalTest(current.state))
            return this.getActions(current);
        let edges = this.stateActionMapping(current.state)
        for (let i=0;i<edges.length;i++)
        {
            let newState = this.transitionModel(current.state, edges[i]);
            let duplicate = false;
            for (let i=0,len=this.closedList.length;i<len;i++)
                if (this.stateEquals(this.closedList[i], newState))
                {
                    duplicate = true;
                    break;
                }
            if (!duplicate)
            {
                let newNode = new TreeNode<T, A>(newState, edges[i], current);
                this.openList.push(newNode);
                this.closedList.push(newState);
            }
        }
    }
    return null;
}
```

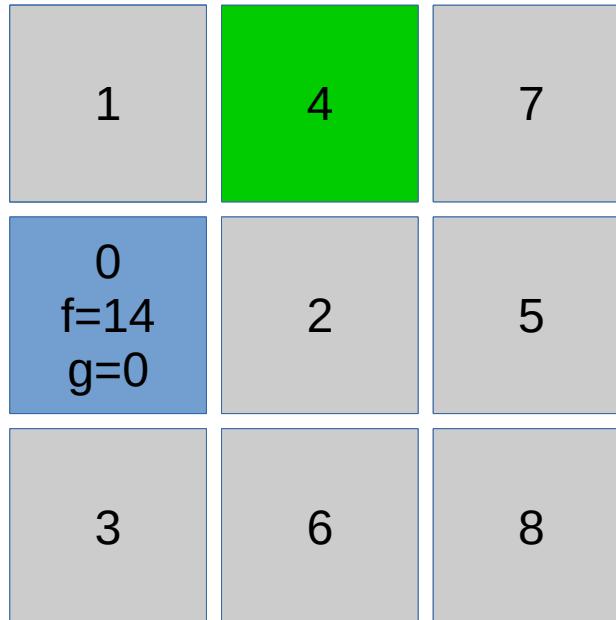
- Pros:
 - Finds shortest path (by number of nodes)
- Cons
 - Unguided

A* Search

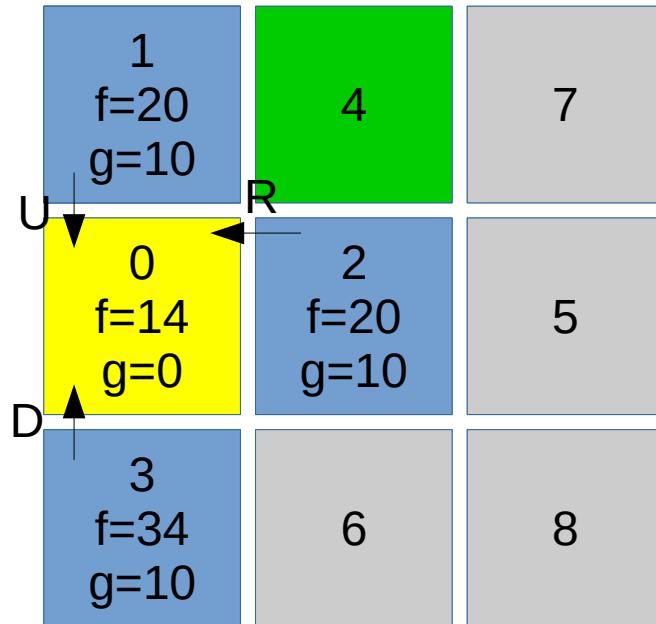


- A* Search
 - Guided / Best-first search
 - Prioritise nodes to explore based on cost + heuristic
 - $g(n)$: cost to reach node
 - $h(n)$: heuristic for node
 - $f(n) = g(n) + h(n)$: priority of node

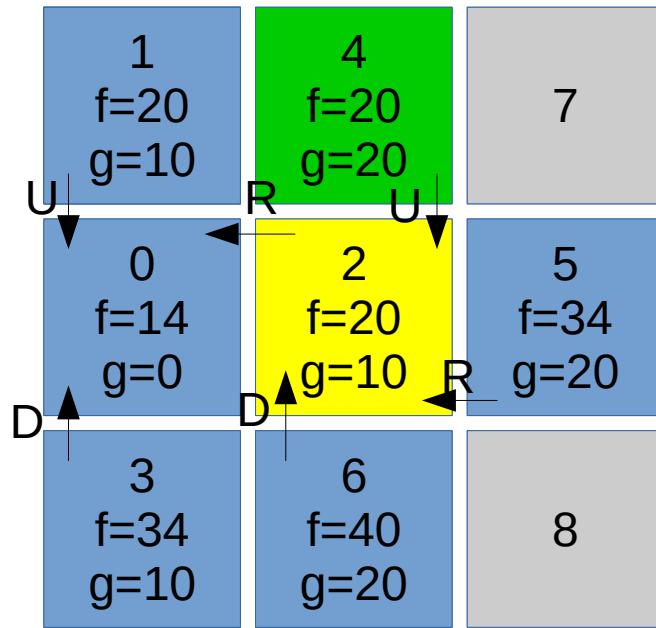
- Set up:
 1. Create empty list (Open set)
 2. Create maps from node → cameFrom, gScore, and fScore (or store $f(n)$, $g(n)$, and parent for each node)
 3. Add start node to Open Set



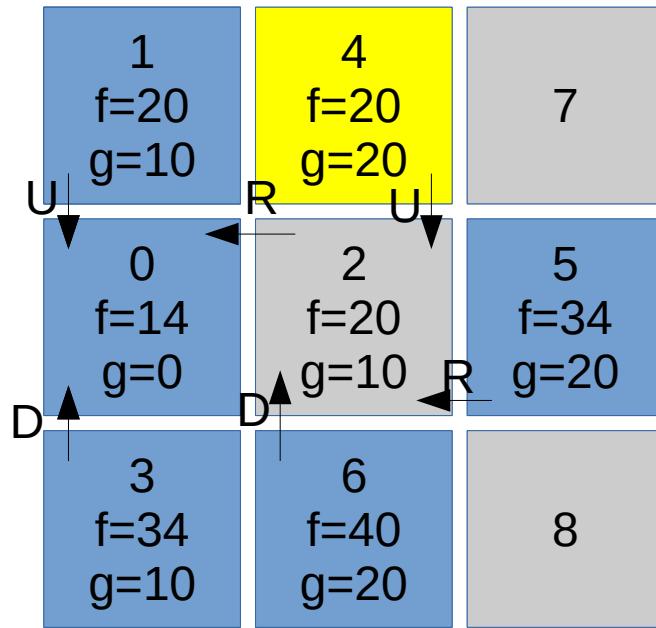
- While open set is not empty:
 1. Remove node in Open Set with lowest $f(n)$
 2. Check if goal
 3. For each neighbour n :
 1. If new $g(n)$ smaller:
 1. Update $g(n), f(n)$
 2. Set parent to current node



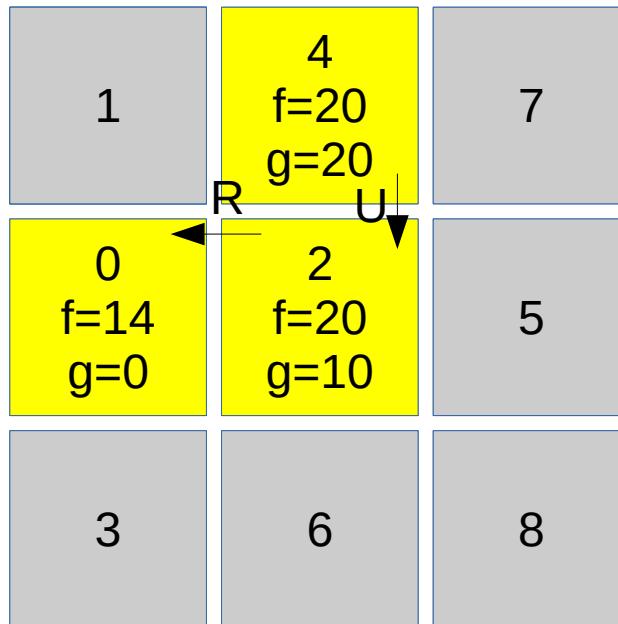
- While open set is not empty:
 1. Remove node in Open Set with lowest $f(n)$
 2. Check if goal
 3. For each neighbour n:
 1. If new $g(n)$ smaller:
 1. Update $g(n), f(n)$
 2. Set parent to current node



- While open set is not empty:
 1. Remove node in Open Set with lowest $f(n)$
 2. Check if goal
 3. For each neighbour n:
 1. If new $g(n)$ smaller:
 1. Update $g(n), f(n)$
 2. Set parent to current node



- Reconstruct path
 - 1. Create empty list
 - 2. Current node = goal node
 - 3. Add action to reach current node to list
 - 4. While current node has parent
 - 1. Set current node = parent
 - 2. Prepend action to react current node to list



Heuristics



- Heuristics
 - **Quick to compute** approximation of value/distance to goal of a given node
 - For A* must be equal to or an under-estimate of true cost
 - Weighted heuristics: $f(n) = w * h(n) + (w-1) * g(n)$
 - If $w = 0$, A* is equivalent to **Dijkstra's algorithm**
 - If $w = 1$, A* is equivalent to **greedy best-first search**

- Manhattan distance

```
function heuristic(node)  
    dx = abs(node.x - goal.x)  
    dy = abs(node.y - goal.y)  
return D * (dx + dy)
```

- Best for 4-connected grids



- Diagonal distance

```
function heuristic(node)  
    dx = abs(node.x - goal.x)  
    dy = abs(node.y - goal.y)  
return D * (dx + dy) +  
        (D2 - 2 * D) *  
        min(dx, dy)
```

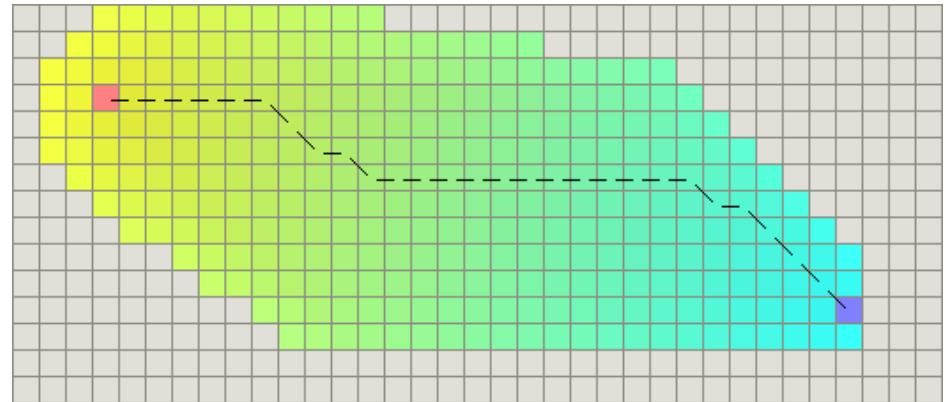
- Best for 8-connected grids



- Straight line distance

```
function heuristic(node)  
    dx = abs(node.x - goal.x)  
    dy = abs(node.y - goal.y)  
    return D * sqrt(dx * dx +  
                    dy * dy)
```

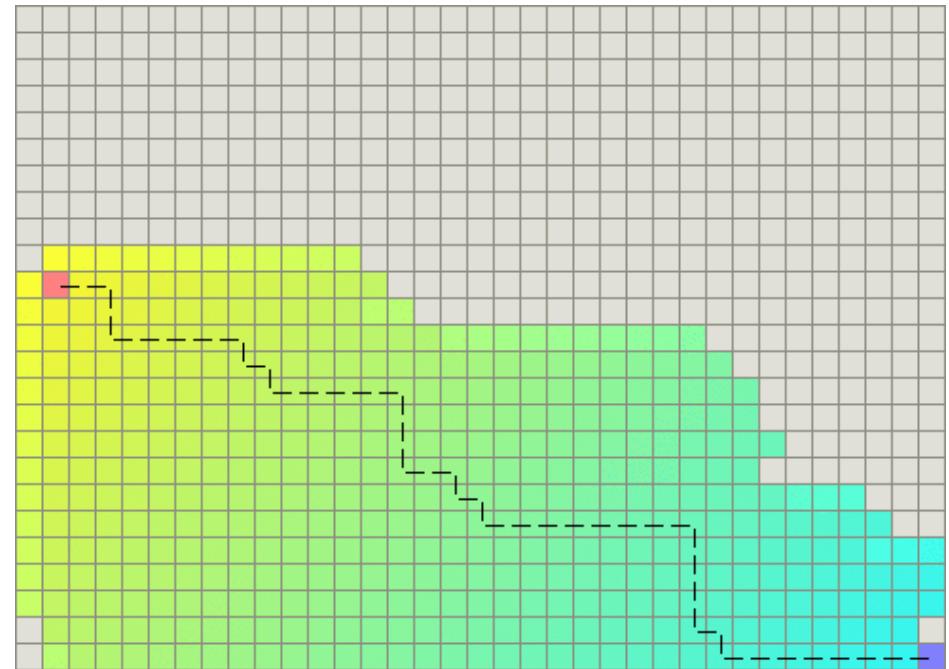
- If you can move at any angle, but leads to more exploration as cost function g and heuristic h will not match



- If you have multiple goals

```
function heuristic(node)  
    h1 = heuristicToA(node)  
    h2 = heuristicToB(node)  
    ...  
return min(h1, h2, ...)
```

- Breaking ties
 - Often many paths with same length, leading to lots of unnecessary exploration
 - Tweaking A* heuristics can improve performance, but impacts on behaviour

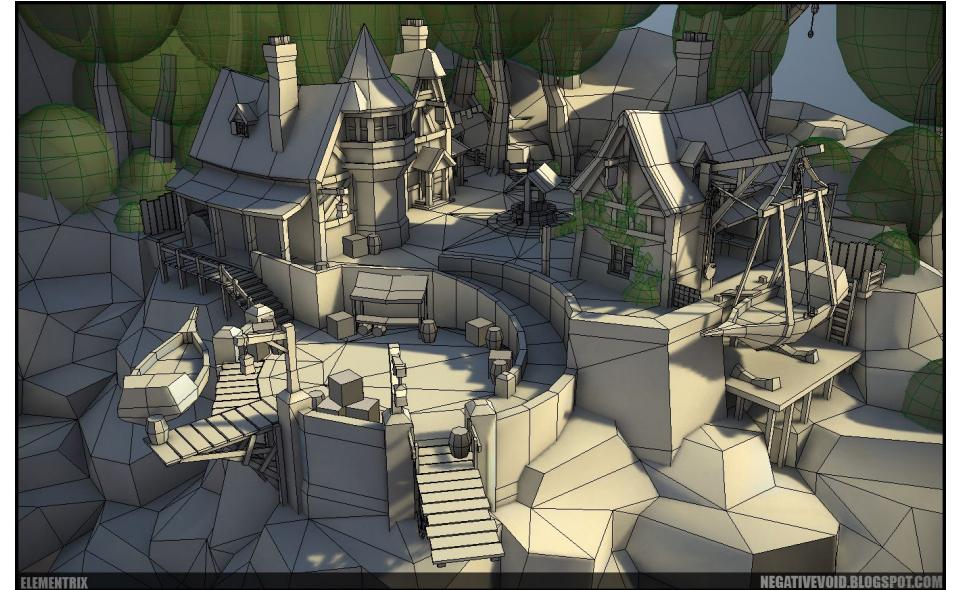


1. Implement an agent that uses A* to plan a path to a goal (e.g. the player)
 1. First **localise** the start and goal positions in terms of grid coordinates
 2. Then **search** to determine a path (sequence of grid cells/directions)
 3. Then **position** the grid cells back in world space and move between them.
2. Explore different options for tie-breaking and see how they affect behaviour and performance:
 - <https://theory.stanford.edu/~amitp/GameProgramming/Heuristics.html>

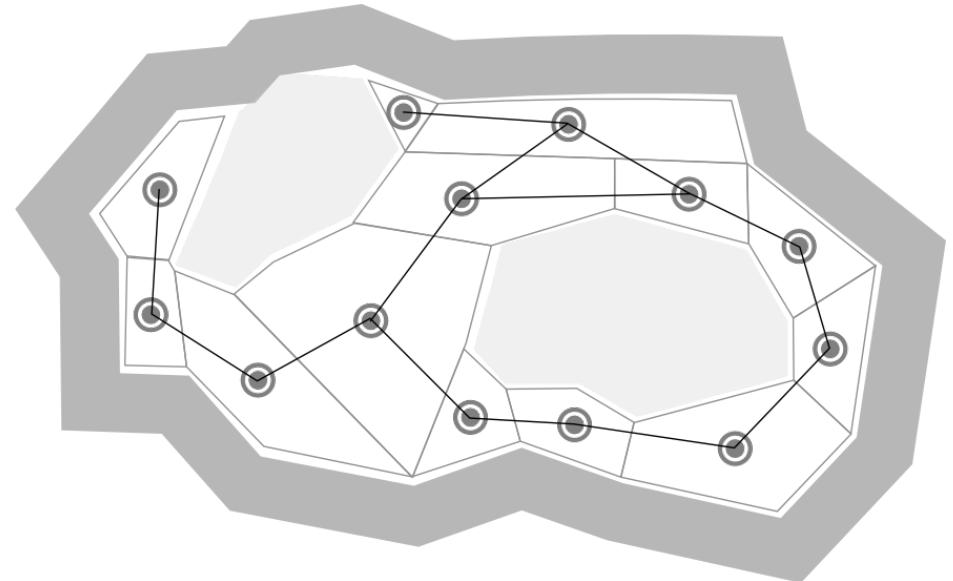
Map Abstractions



- Game worlds are usually unsuitable for navigation
 - “Polygon soup”
 - Doesn’t encode navigable areas
 - Irrelevant detail



- Navigation graph
 - Graph size affects
 - Search performance
 - Path quality
 - Content creation problem



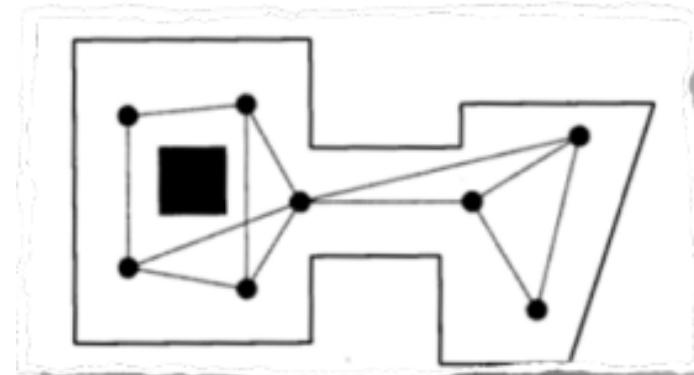
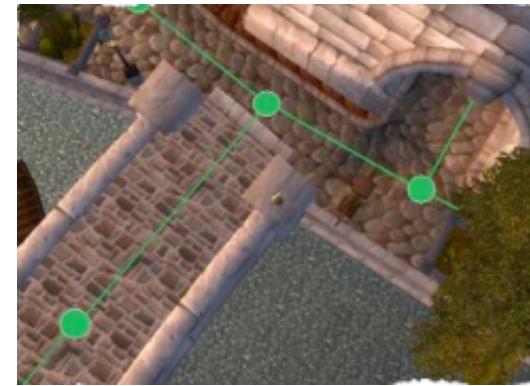
- Pathfinding tasks
 - **Localisation**: convert position to node
 - **Search**: find node path in nav graph
 - **Positioning**: convert node(s) to position(s)
 - **Smoothing**: improve quality of path

- Requirements
 - **Validity:** If two nodes are linked they are navigable in the world (the graph doesn't mislead)
 - Invalid structures result in paths that agent's can't follow
 - **Completeness:** All navigable spaces in the world are represented in the graph (the graph doesn't omit)
 - Incomplete structures prevent agents taking valid paths

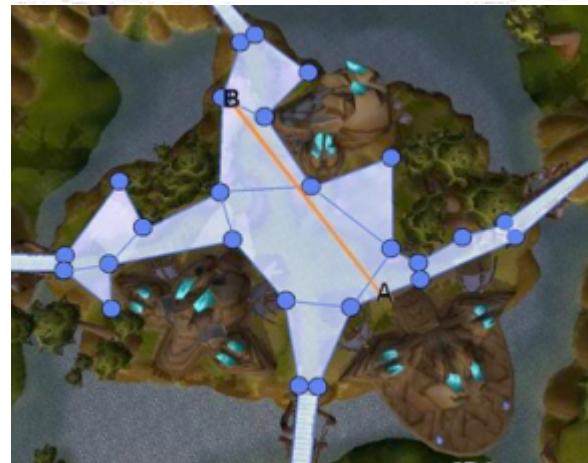
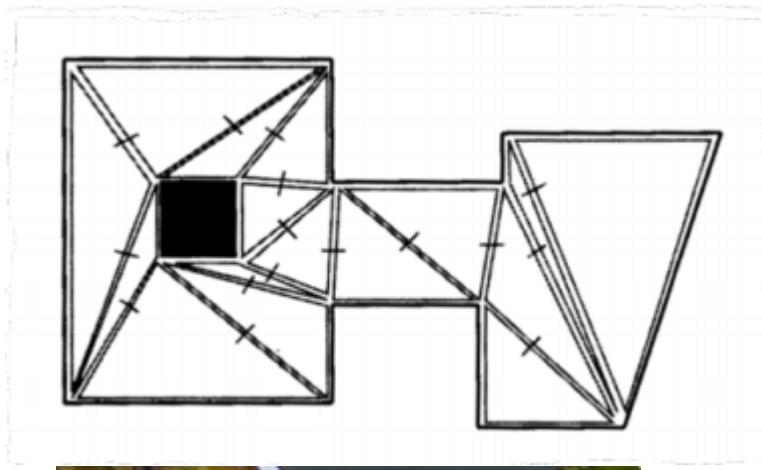
- **Grid Navigation**
 - Divide world into regular regions
 - Easy to generate
 - Large numbers of (unnecessary?) nodes



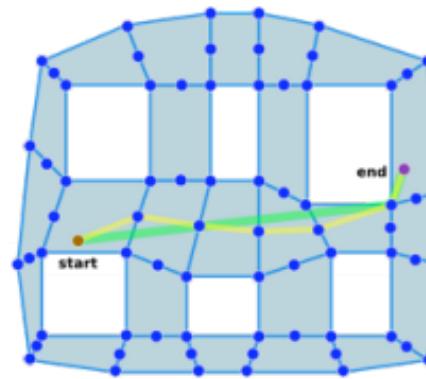
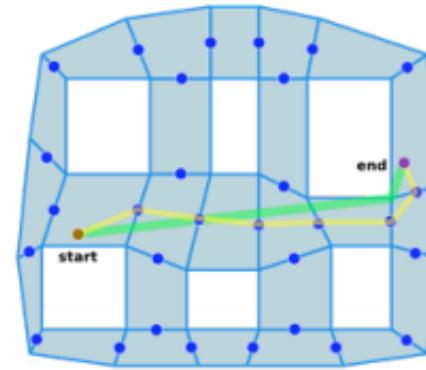
- **Waypoint Navigation**
 - Place waypoints and connections in the world
 - Difficult to generate
 - Localisation searches for nearest waypoint, but we need to validate movement to the waypoint (raycast)



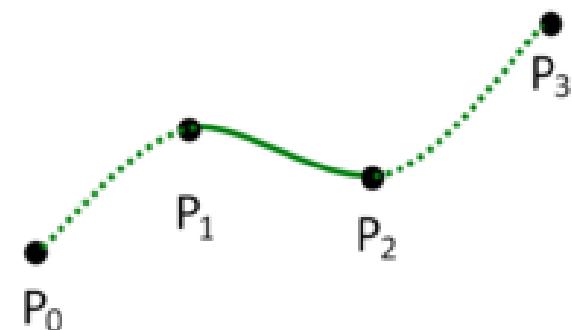
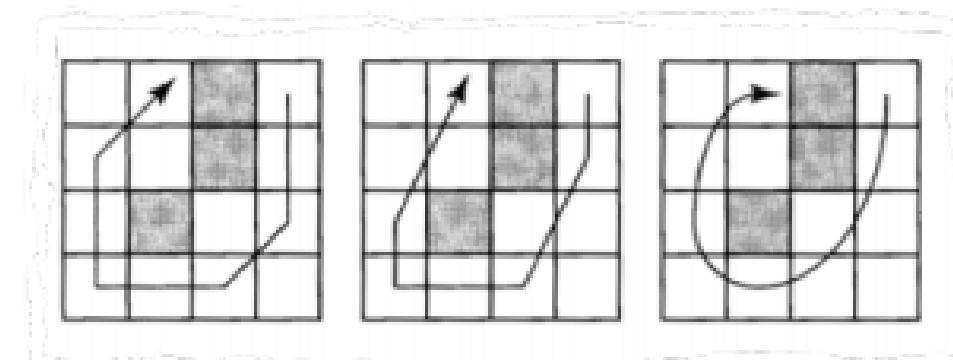
- **Navigation Meshes**
 - Set of polygons (often triangles) that describe navigable surface
 - Better paths from smaller graphs
 - Can generate from level geometry



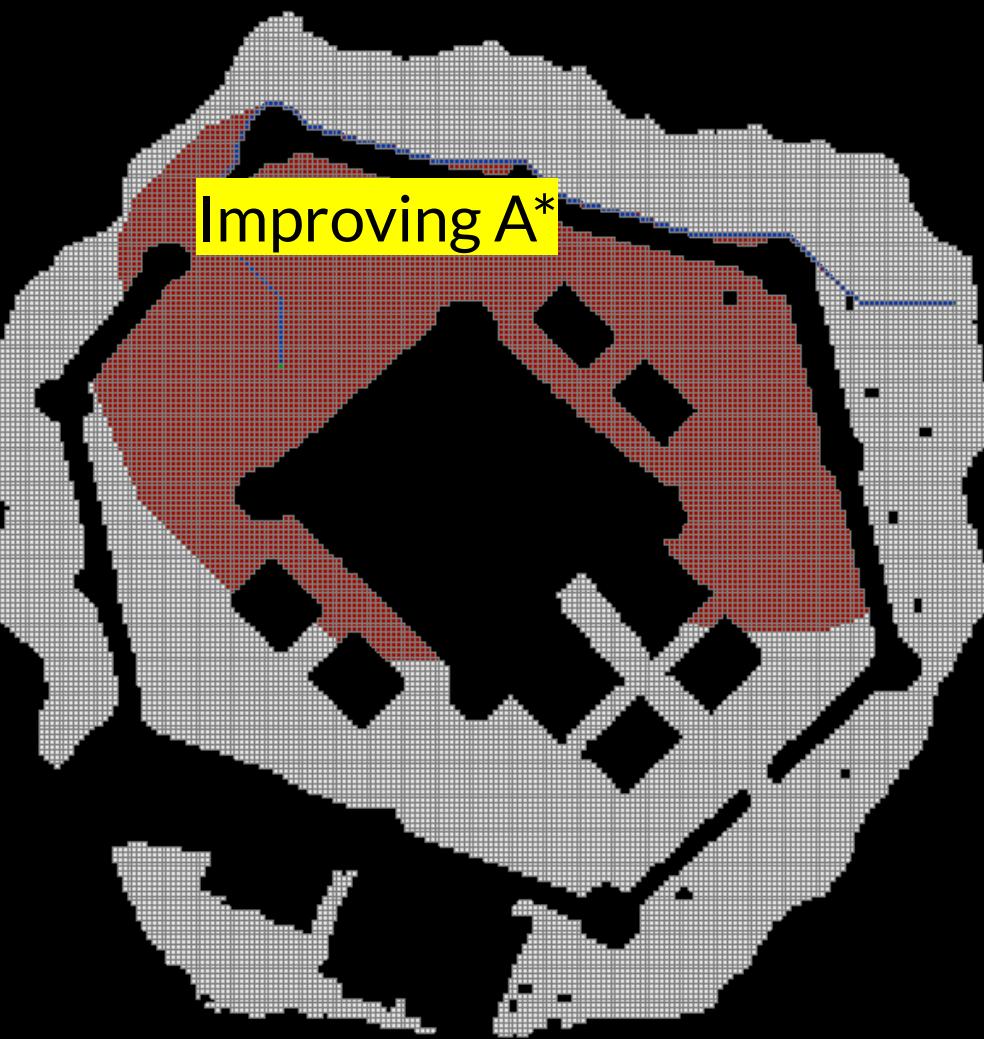
- Navigation Meshes: **Pathfinding**
 - Pathfinding returns a series of polygons
 - Polygons can be crossed using edge midpoints.
 - Adding vertex points improves cornering
 - The path remains valid as long as the same polygons are used



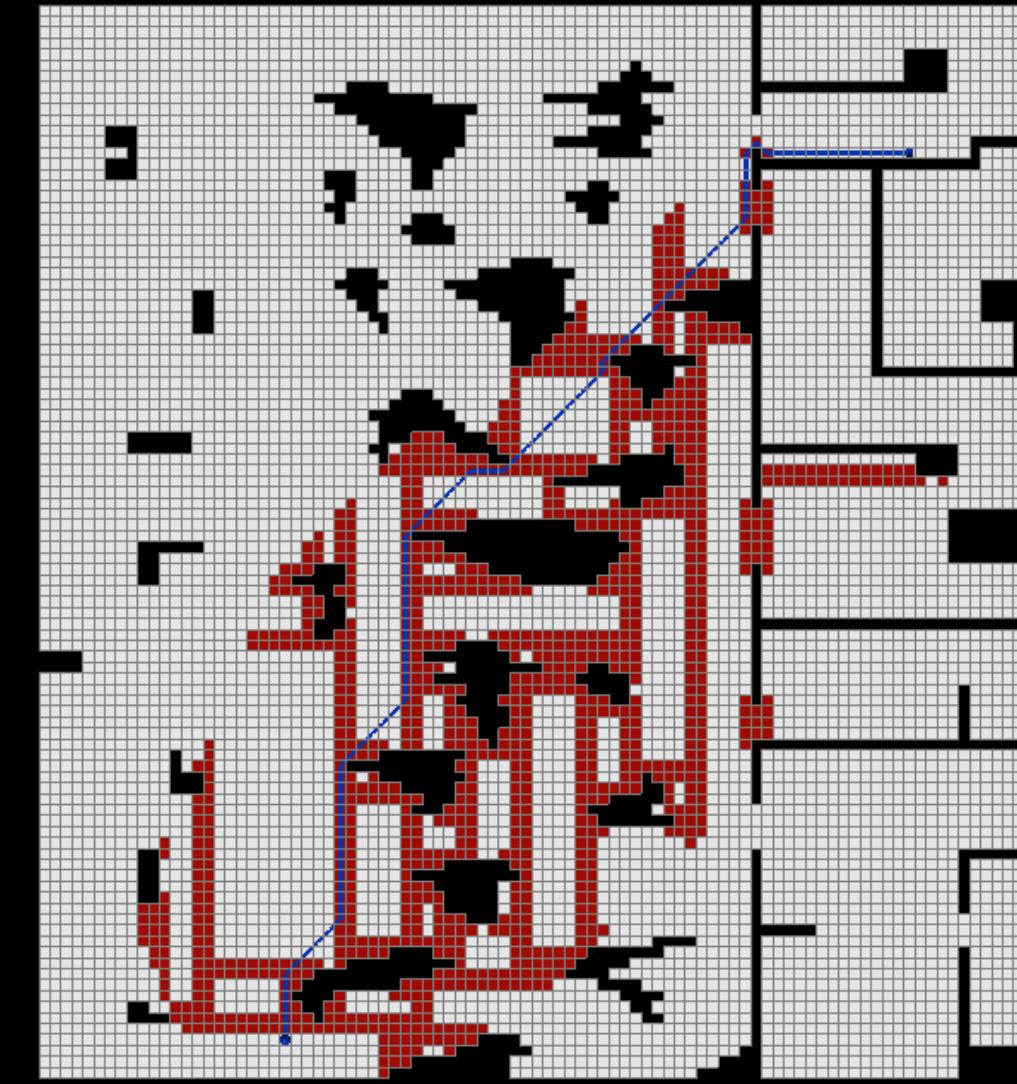
- **Path smoothing**
 - **String pulling** checks three consecutive points at a time: if the line from the first to the third is navigable, then the second is removed
 - A smooth path that pass through the remaining points can be computed, e.g. using Catmull-Rom splines



SIMULATION TIME ELAPSED: 9.14, DISPLAY TIME: 9.14



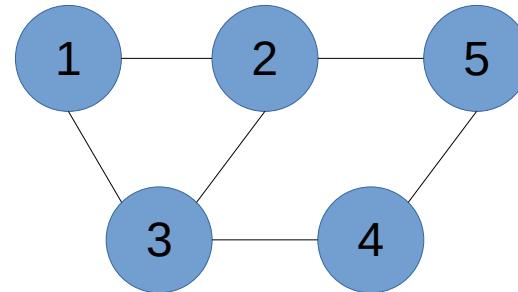
SIMULATION TIME ELAPSED: 10.14, DISPLAY TIME: 10.14



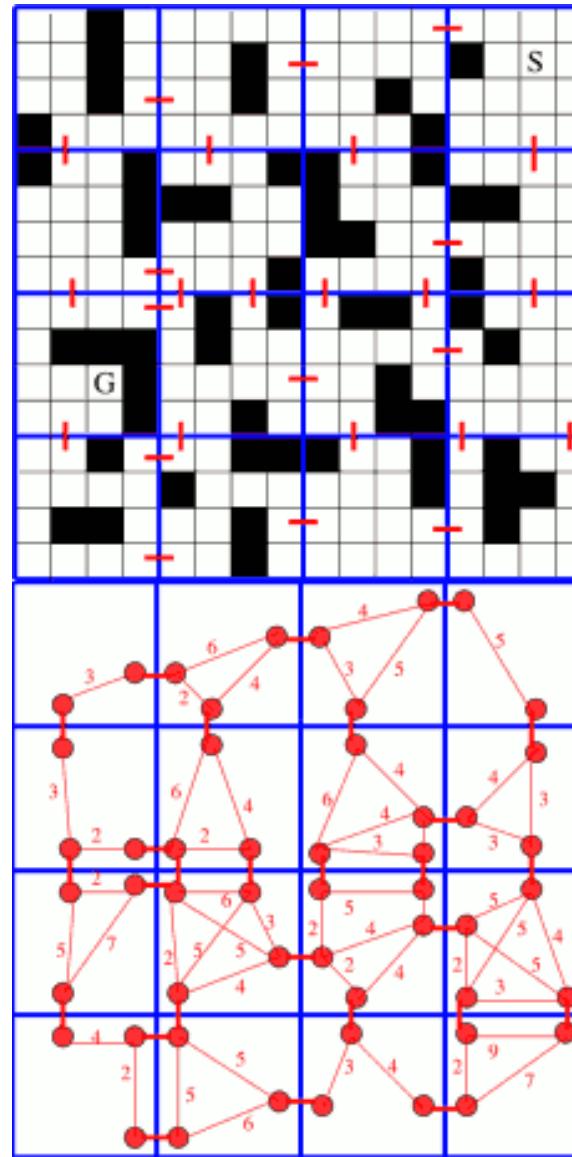
- **Memory Heuristics:** Store precomputed distances between key locations
 - Fast
 - Optimal
 - High memory
- **Abstraction:** reduces size of search space
 - Fast
 - Low-memory
 - Suboptimal paths
- **Search-space pruning:** Identify nodes that don't need to be explored
 - Slow
 - Low memory

- **Precomputing paths**
 - “Next node” lookup table
 - For given current, goal, return next node/edge in path
 - Time: $O(n)$
 - Space: $O(n^2)$
 - Alternatively store as edge number as there are usually fewer edges per node than nodes

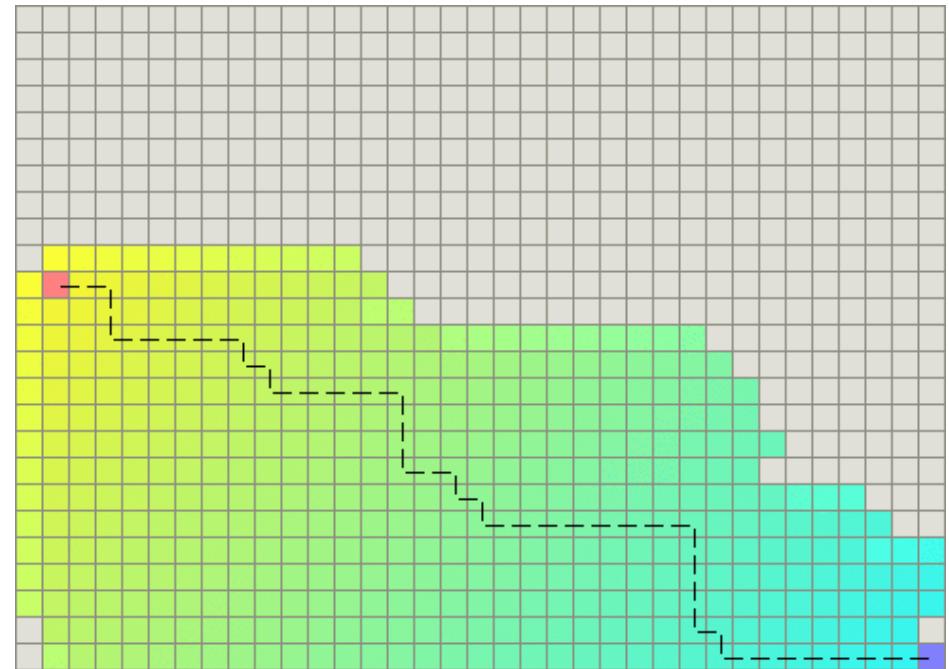
Current node	1	2	3	4	5	Goal node
1	-	2	3	3	2	
2	1	-	3	3	5	
3	1	2	-	4	2	
4	3	3	3	-	5	
5	2	2	2	4	-	



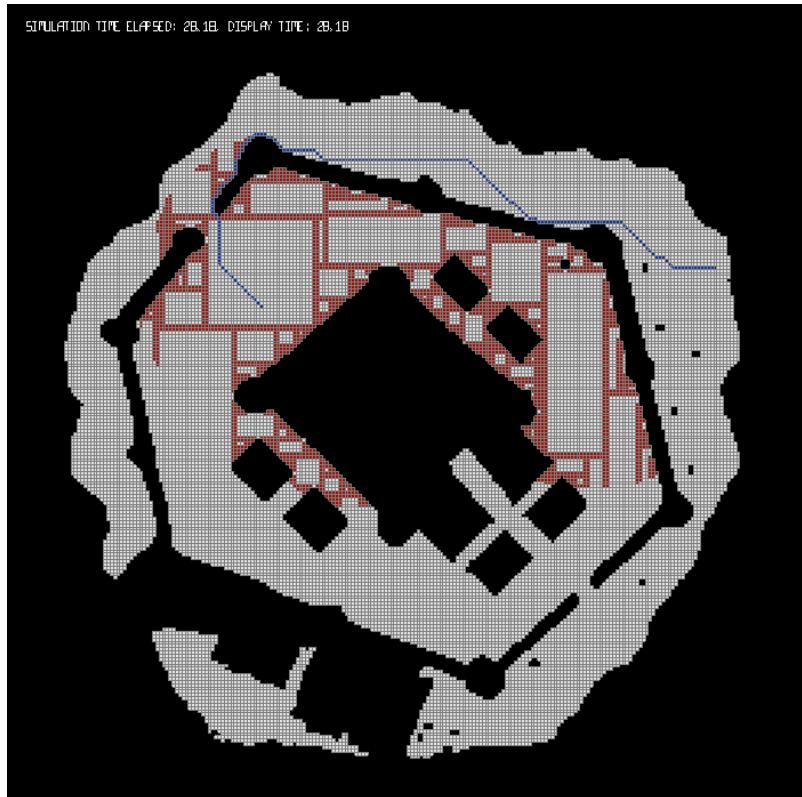
- **HPA***: Hierarchical variant of A*
 - Pre-computed abstract graph
 - Group low-level nodes into clusters
 - Each cluster is represented by a subset that borders other clusters
 - A* from start and goal to members of this cluster
 - A* over abstract graph
 - Refine into a path on original graph using A* or cached paths



- **Symmetry breaking**
 - In grids, you often get paths that only differ in the order of moves
 - Performance gain from avoiding unnecessary work



- Rectangular Symmetry Reduction (RSR)
 - Decompose map into empty rectangles
 - Only expand nodes on perimeter of rectangle



Boulder's Gate A* + RSR

- Jump Point Search (JPS)
 - Avoid expanding neighbours that could be reached more optimally by parent
 - Make jumps across open spaces
 - More expensive per node, but fewer nodes



Boulder's Gate A* + JPS

- **Dynamic pathfinding**, e.g. D* Lite
 - Plan in unexplored space or where path costs are changing
- **Angle-angle pathfinding**, e.g. BlockA*
 - Generates shortest straight-line paths using line-of-sight checks
 - Path smoothing as part of the algorithm
- **Incremental pathfinding**, e.g. LPA*
 - Re-uses information generated for previous (similar) searches
- **Low memory use**, e.g. IDA*

1. Experiment with improvements to A* in your pathfinding agent
 - Breaking ties
 - <https://theory.stanford.edu/~amitp/GameProgramming/Heuristics.html>
 - Path smoothing
 - Precomputing paths
 - Symmetry Reduction
 - <https://harablog.wordpress.com/2011/08/26/fast-pathfinding-via-symmetry-breaking/>
 - Implement an A* variant algorithm

- More resources
 - Amit's A* Pages
 - <http://theory.stanford.edu/~amitp/GameProgramming/>
 - The Moving AI Lab @ U. Denver movingai.com