



Income:
£124.66 Bn

Expenditure:
-£129.55 Bn

Deficit:
£4.90 Bn

Debt:
£450.47 Bn



Games AI

Lecture 3.1

Utility-based AI

- Utility-based AI
 - Utility theory
 - Decision factors
 - Designing response curves
 - Implementing utility-based AI

- Utility theory
 - General model of **decision making**
 - **Every possible action** can be scored on a single scale of “goodness” (utility, weight, score, priority, ...)
 - Agents act rationally to **maximise utility**
“Irrationality” / cognitive biases / personality traits can be added into the model as factors

- Utility != value
 - Value is a measurable quantity, eg. price
 - Utility is how much you desire something
 - Money has fixed value but variable utility depending on context
 - eg. £20 on a desert island is worthless
 - Utility is abstract, but it can be based on values

- Principle of Maximum Expected Utility
 - Multiply utility of each outcome (D) by probability of it occurring (P)
 - e.g. attack in D&D
 - Critical hit: 5%, utility: 0.9
= 0.045
 - Regular hit: 90%, utility: 0.6
= 0.54
 - Fumble: 5%, utility:-1
= -0.05
- ∴ expected utility of attack = 0.555

$$EU = \sum_{i=1}^n D_i P_i$$



- Utility-based AI basics
 - Every action is scored at once using a mathematical model of the value (utility) of each action
 - Top scoring actions change as scores fluctuate dynamically
 - One of top scoring actions is chosen

- Pros
 - More nuanced behaviours than using booleans
 - Hand authored → Tight authorial control
 - Simple to implement
 - Considerations are reusable between agents
 - Considerations reflect decision making process - can be exposed to the player
- Cons
 - Designing behaviours can be a “dark art”
 - No planning or look-ahead



Making Decisions



\$580,400
941 Guests

RollerCoaster Tycoon
© Copyright 2000 Chris Sawyer

March, Year 3
59F

- Decision factors
 - Multiple pieces of data weigh into any given decision
 - Health, Ammo, Distance to enemy
 - Wealth, hunger, enjoyment of rollercoasters
 - Conversation choices, loyalty, distance to allies
 - Almost anything can be a decision factor
 - Cooldown, repeat penalty, aggressiveness, conscientiousness, ...

Making decisions

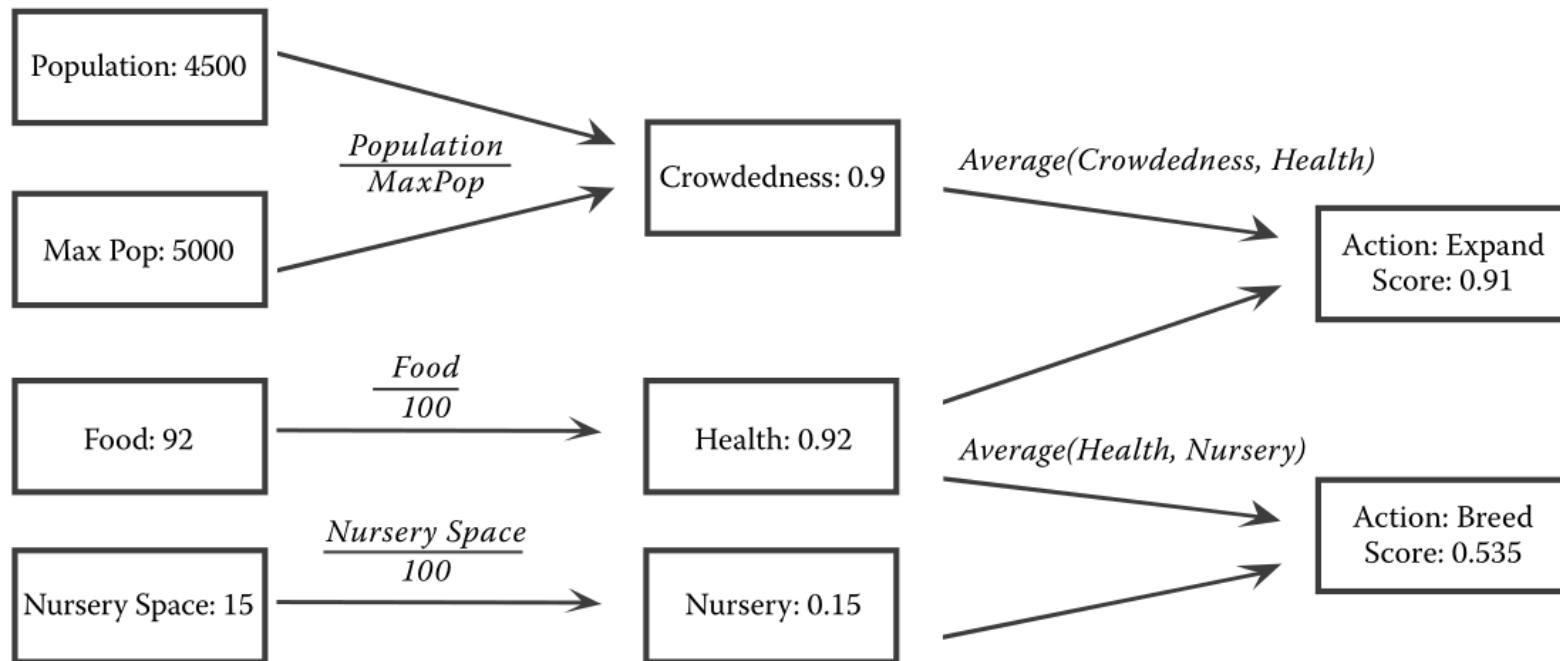
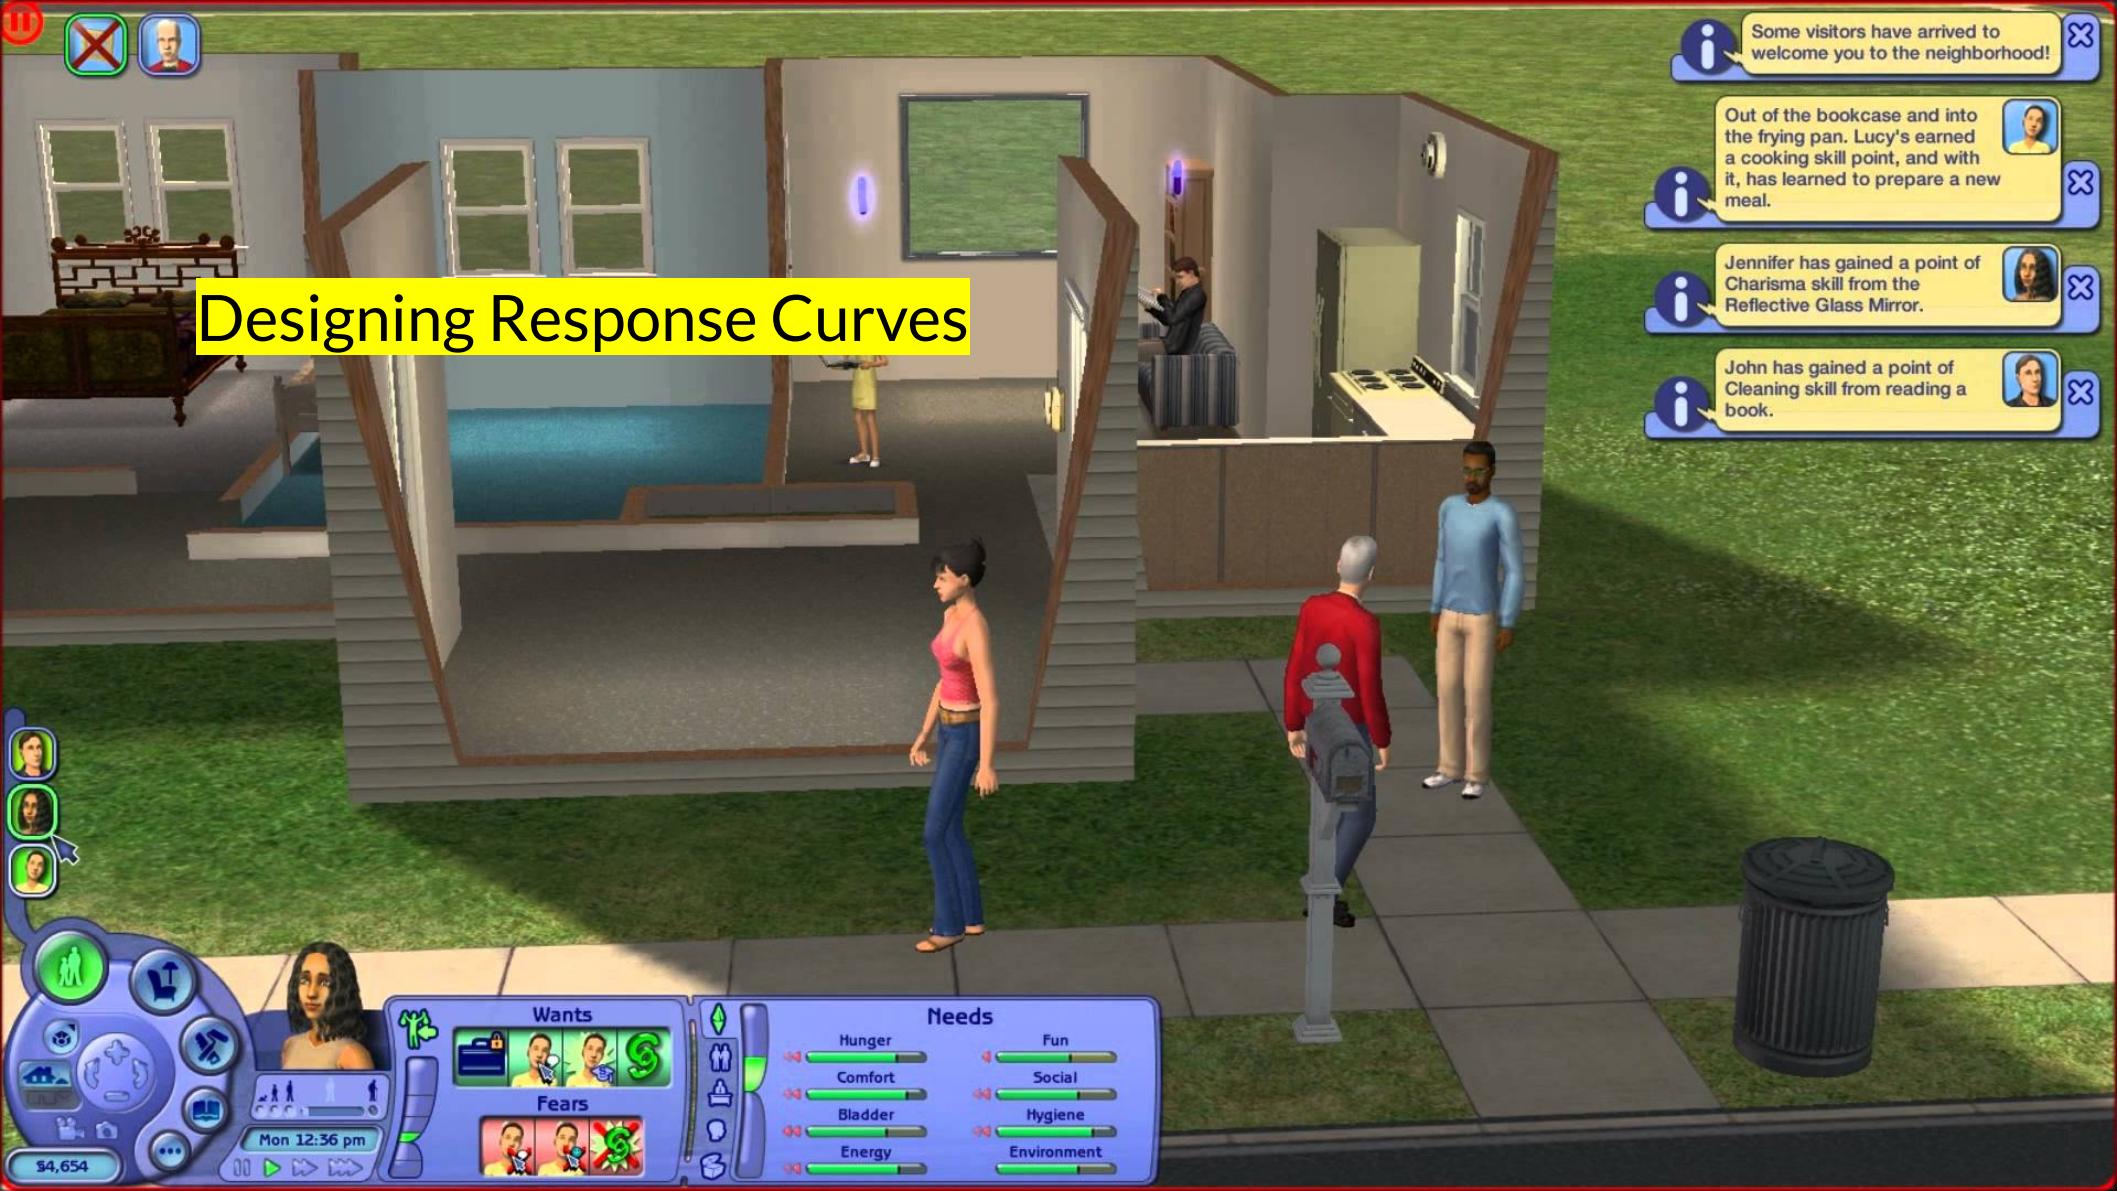


Figure 9.1

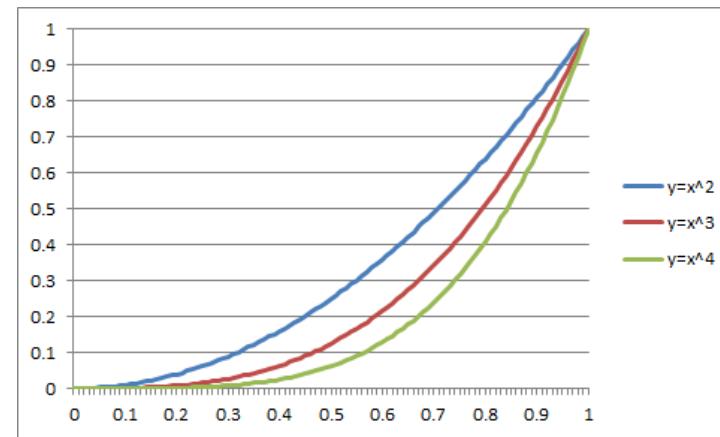
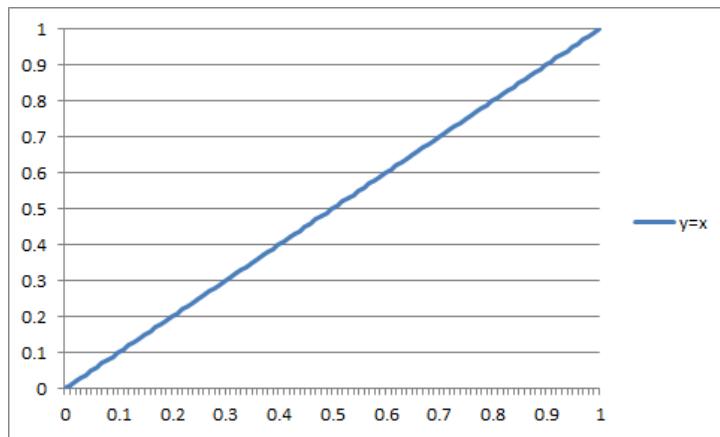
An example of combining utility scores from different decision factors to arrive at a final score.

- Take with highest score (absolute utility)
 - Take “best” decision
 - Rigid, predictable
- Select an option at random, but weight probabilities (relative utility)
 - Variation, interesting
 - Periodically will do something that makes your AI look stupid
- Combine the two for the best results

- Decision factors
 - Art of design is mostly selecting considerations to weight actions and pairing with appropriate response curves
 - By averaging normalised scores can build endless chain of combinariions



- What is a response curve?
 - A mathematical representation of a relationship modelling how a change in the input affects the output

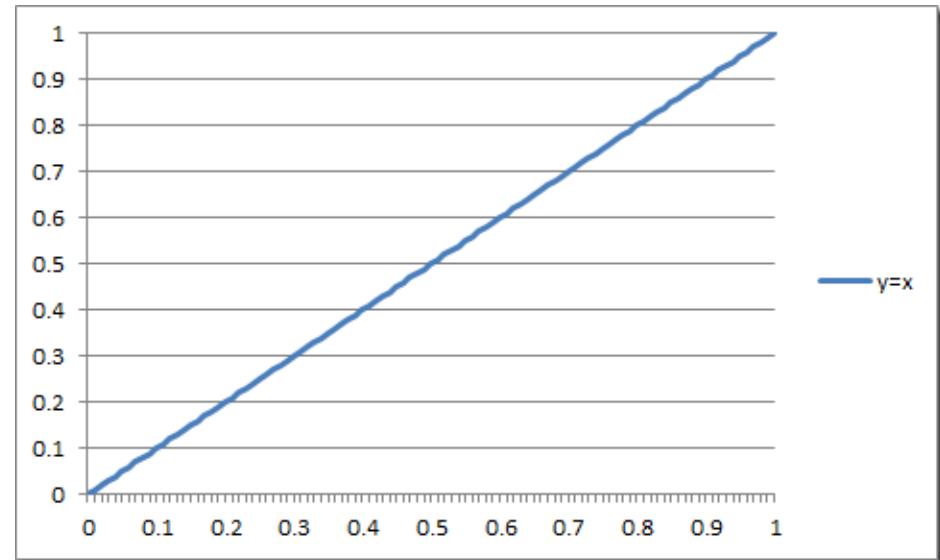


- A response curve is a function with domain and range in $[0-1]^*$
 - It maps some normalised input value (eg. health %) to a consideration for our AI
 - eg. health% → desireToHeal
 - eg. ammo% → desireToReload
 - Useful properties:
 - Multiplying curves always gives result in $[0-1]$
 - Multiplying by 0 always gives 0
 - Can combine any number of curves
 - Outputs of one curve can be inputs to another

*or another consistent range, but mathematical properties may change

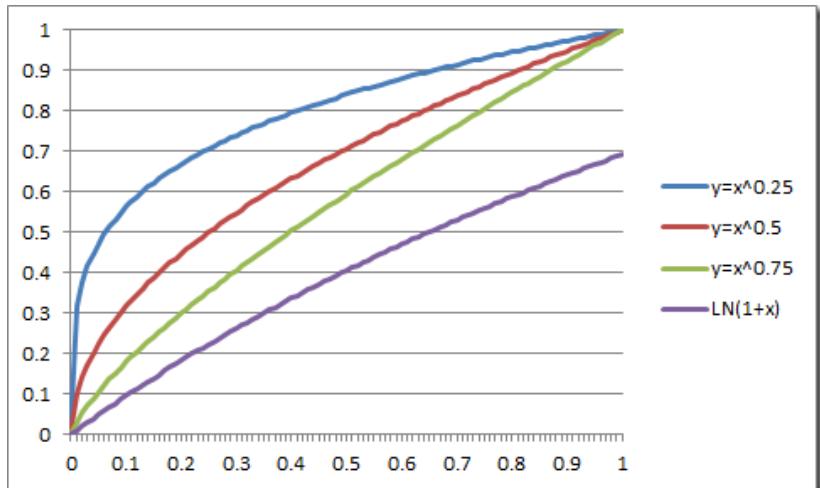
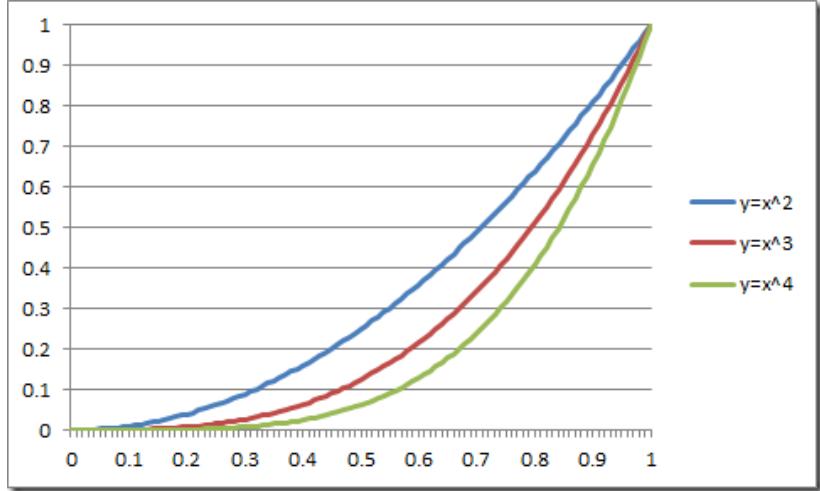
- Four types of curve
 - Linear/Quadratic
 - Exponential
 - Logistic
 - Logit

- Linear
 - $y=mx + c$
 - Change in x always causes the same change in y

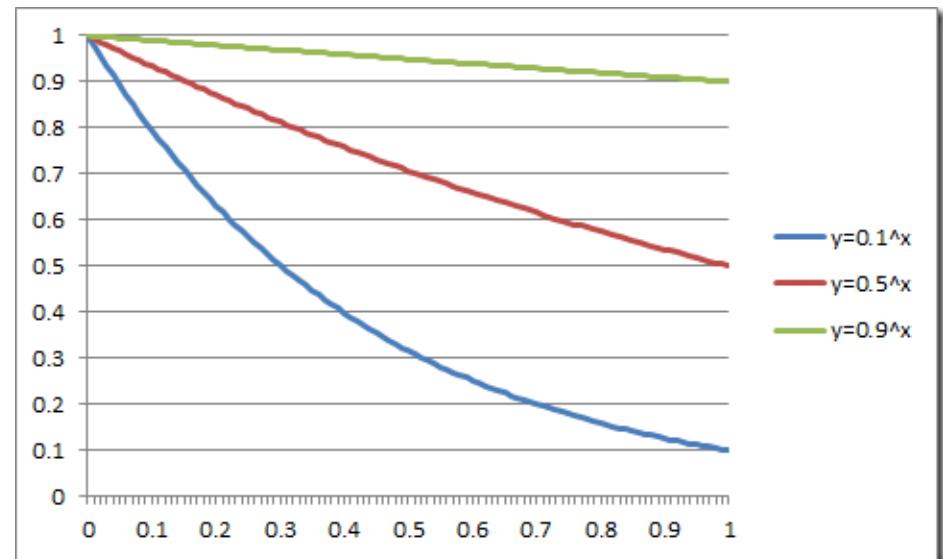


Designing response curves

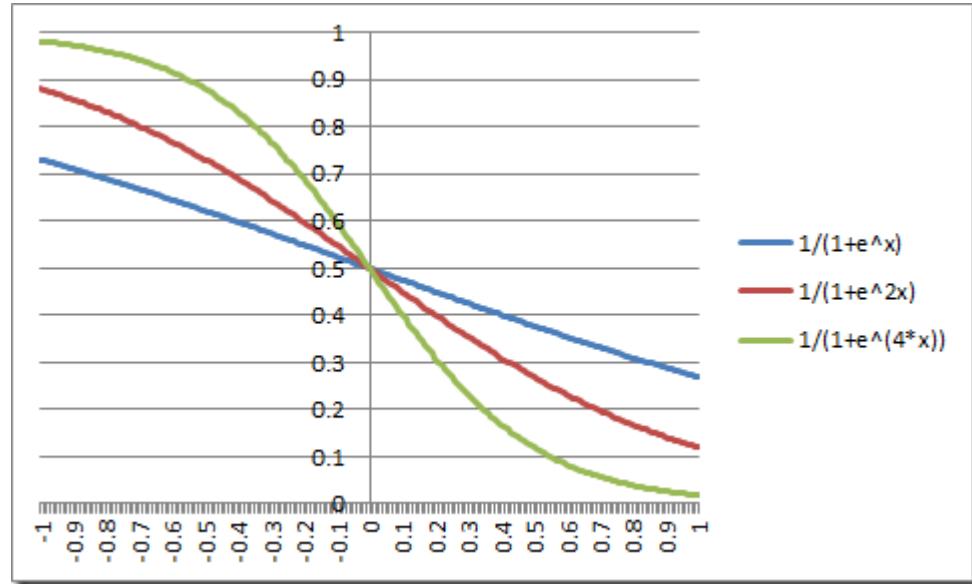
- Quadratic
 - $y=x^a$
 - $a > 1$ increasing gradient
 - $0 < a < 1$ decreasing gradient
 - Changes in x lead to bigger/smaller changes in y



- Exponential
 - $y=a^x - 1$: exponential growth
 - $y=a^x$, where $0 < a < 1$: exponential decay

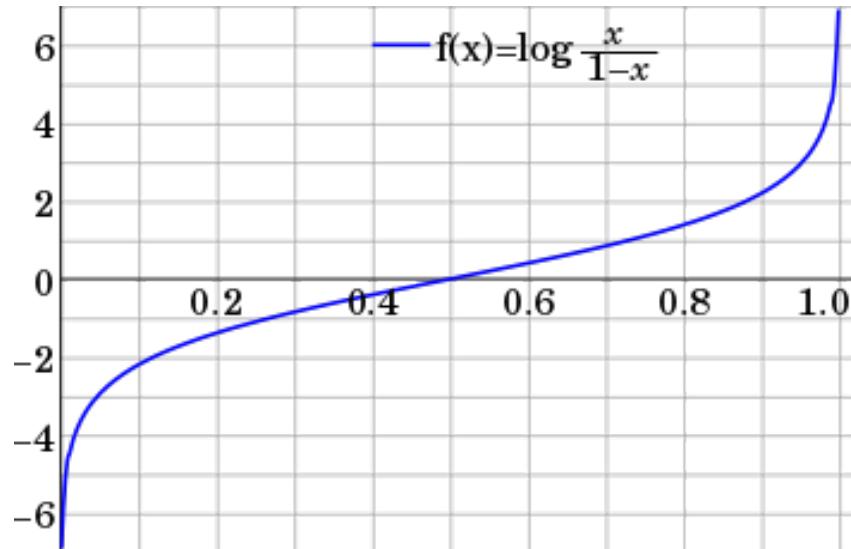


- Logistic
 - $y = 1/(1+e^x)$
 - $y = 1/(1+e^{-x})$ for reverse
 - “Soft threshold” giving a more or less sudden change between low and high



<https://www.desmos.com/calculator/j7ucn0rkfc>

- Logit curve
 - $y = \log(x / 1-x)$
 - Differences at the extremes are more significant than those in the middle



<https://www.desmos.com/calculator/q9hkfpqc2s>

- Can customise your own functions using these mathematical components
 - Need to adjust to fit within your scale (e.g. 0-1)
 - Clamp if necessary
- Much easier:
 - Paramatise these four basic function types

- These four curve type can all be paramatised: b,c,k,m
 - Linear/Quadratic
 - $f(x) = m(x-c)^k + b$
 - Exponential
 - $f(x)$
 - Sigmoid/Logistic
 - $f(x) = (k/(1+1000e^{m^{-1}x+c})) + b$
 - Logit
 - $f(x) = 0.5 \log_{100^m} ((x/k - c)/1-(x/k - c)) + b + ...?$

- Linear/Quadratic
 - $f(x) = m(x-c)^k + b$
 - m = slope
 - k = how much bend
 - b = vertical shift
 - c = horizontal shift

- Logistic
 - $f(x) = (k/(1+1000em^{-1x+c})) + b$
 - m = slope at inflection point
 - k = vertical size of curve
 - b = vertical shift
 - c = horizontal shift



- Each response curve can be defined by an enum and four floats:
 - [{Quadratic, Exponential, Logistic, Logit}, m, k ,b ,c]

Curve type	m	k	b	c
Quadriatic	0.5	0.1	0.25	0.0
Logistic	50,	-0.95	1.0	0.6

- Processing a response curve
 - Function takes parameterised curve and input x
 - $(type, m, k, b, c, x)$
 - Clamp input to $[0-1]$
 - Apply function determined by curve type
 - Clamp output to $[0-1]$
 - Return y

- Processing a consideration (decision factor)
 - Each consideration has an array of inputs
 - Inputs can come from other considerations or from a “Clearing House” function that returns normalised input values (eg. health, ammo)
 - Each input is assigned a paramatised utility function
 - The outputs of these functions are multiplied together
 - The result is the utility of the consideration

- Processing an action
 - Each action has an array of considerations
 - The utility of each consideration is calculated and multiplied together
 - The result is then multiplied by a weight, e.g.
 - Idle actions: x1
 - Attack actions: x3
 - Scripted actions: x5
 - Combat responses: x10

- Calculate the score of each action
 - Score actions with targets on a per-target basis
- Push scored actions onto a list of possible actions
- Pick (one of the) top scoring actions from the list



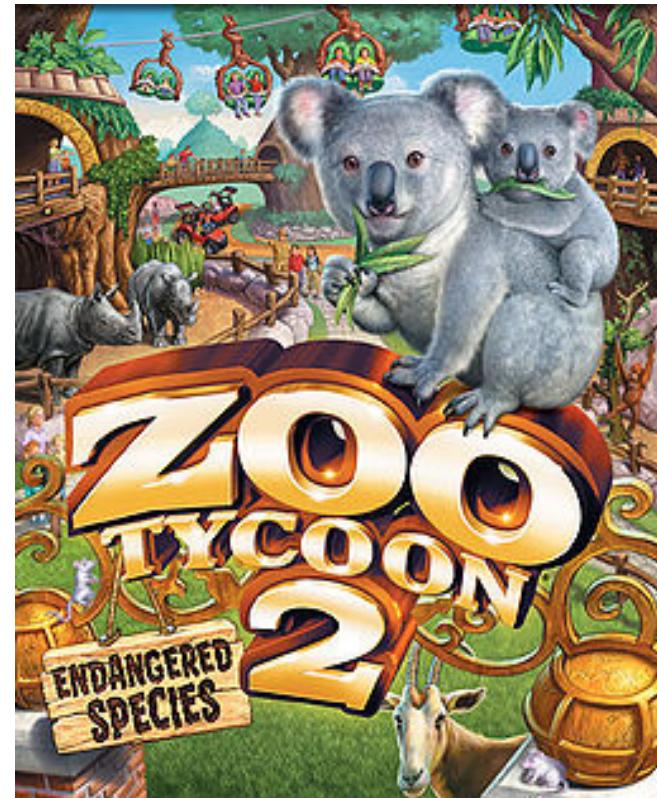
Implementing Utility-based AI (2)

Dual-utility Reasoning

- Each consideration has
 - A utility
 - A rank
- Rank of action is maximum rank of any consideration
 - i.e. considerations can force overall action to be of high rank
- Weight of action determined by the product of utility of considerations

- Dual-utility reasoning
 - 1. Eliminate the options with a weight of zero.
 - 2. Find the highest rank category, and eliminate options that don't belong to it.
 - 3. Find the best remaining option (i.e., the one with the highest weight), and eliminate options that are much worse (i.e., much lower weight) than it.
 - 4. Use weight-based random to select from the options that remain
- Kevin Dill. *Dual-Utility Reasoning* in Game AI Pro 2, ch 3

- Dual-utility reasoning eg.
 - Normal needs rank 0
 - Hunger, entertainment, bathroom
 - Context-specific behaviours higher ranks
 - e.g. only tree climbing behaviours when climbing trees
 - Death rank of 1,000,000



Further “reading”

- David “Rez” Graham. *An Introduction to Utility Theory* in Game AI Pro. Ch.9
 - <http://www.gameaipro.com/>
 - Also several other chapters in the Game AI Pro books about utility systems
- Dave Mark’s Infinite Axis Utility System Architecture at GDC
 - <http://www.gdcvault.com/play/1018040/Architecture-Tricks-Managing-Behaviors-in> (third speaker)
- ...and more of his GDC talks about Utility Systems:
 - <http://intrinsicalgorithm.com/IAonAI/2013/02/both-my-gdc-lectures-on-utility-theory-free-on-gdc-vault/>
- Kevin Dill *Dual-Utility Reasoning* in Game AI Pro 2. Ch.3
- www.desmos.com Free online graphing tool