

Games AI
Lecture 6.2




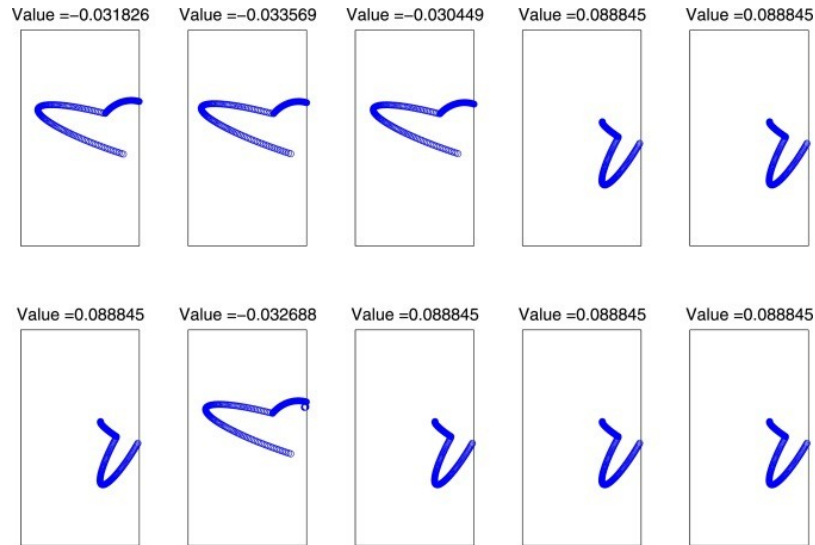
Evolutionary Algorithms




- Last lecture we described PCG as a search problem
 - 1) You have a **generative space** defined by your algorithm
 - 2) Sometimes this space includes both good and bad content.
 - 3) If you can score how “good” content is with an **evaluation function** then you can **search** for high-scoring content.
 - This is an **optimisation** problem

- Why evolution?
 - Engineers often copy solutions in nature
 - What is the best problem solver in nature?
 - **The evolutionary process**


- So you can... find the optimal design for hypersonic re-entry vehicle nosecone



(a)



(b)



(c)

- Or... evolve a more efficient antenna for an unusual purpose



- Or... tune parameters for **rules and FSMs**

```
if (enemy.distance <= 5)
```

```
    attackWithKnife()
```

```
else if (enemy.distance > 5 AND enemy.distance <=30)
```


```
    attackWithSubmachineGun()
```

```
else attackWithRifle()
```

- Or... tune parameters that **define AI personality** (e.g. unit preference, scientific advance preference, offense vs. defense, etc.)
- Ponsen, Marc, and Pieter Spronck. Improving adaptive game AI with evolutionary learning. Diss. Masters Thesis, Delft University of Technology, 2004.
<http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.109.6055&rep=rep1&type=pdf>



- Or... **evolve agents** using e.g neuroevolution (evolution of neural networks)
 - https://youtu.be/_1TOKKgAock




Natural evolution



- Darwinian natural selection:
 - **Competition-based selection**
 - **Phenotypic variation**
 - Behavioural and physical traits that affect an individual's fitness



- Success in survival and reproduction is determined by **phenotypical** properties
- **Phenotypic variations** are always caused by **genotypic variations**
- Phenotypes never influence genetic information
 - **No learning within individuals**



- Change happens through reproduction by the mechanisms of:
 - Mutation
 - Recombination

- **Adaptive landscape**
 - Metaphor of a space where height corresponds to fitness
 - **Global Optimum** = best
 - **Local Optimum** = better than neighbours




Fig. 2.2. Illustration of Wright's adaptive landscape with two traits

- **Each individual is a sample** of the space of all possible living things
 - Produced by forces of **variation**
 - Evaluated by forces of **selection**
 - Prove viable to live
 - Prove viable to reproduce
- In algorithmic terms this is a “**generate and test**” method.

Intro to Evolutionary Algorithms



What?

MAGIKARP is evolving!

- In an Evolutionary Algorithm, a solution is described by a **genotype**
 - This might be a string, enum[], number[]...
 - e.g. **parameter values** for a PCG algorithm
 - e.g. **weights** for a neural network
- This data (genotype) is interpreted to create the **phenotype** (e.g. level map, AI behaviour, game audio)

- Evolutionary Algorithms (EAs) are:
 - **Population based**
 - (You keep hold of a bunch of possible solutions at the same time)
 - **Stochastic**
 - (Involves randomness)
 - Most use **recombination**
 - (You mix together different solutions to try and find a better one)

- Why EAs work
 - **Variation** (recombination and mutation) create diversity and novelty
 - **Selection** acts as a force to increase the mean quality of solutions

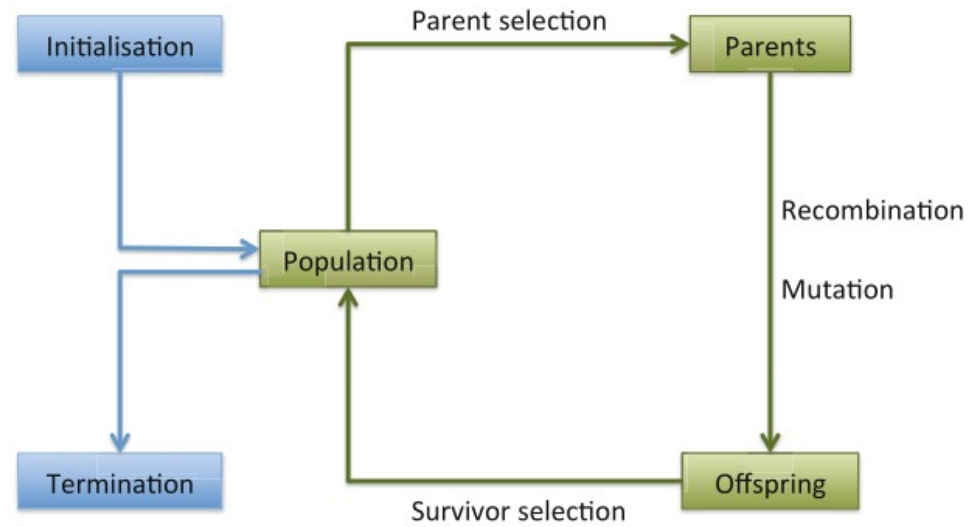


Fig. 3.2. The general scheme of an evolutionary algorithm as a flowchart

```
BEGIN
  INITIALISE population with random candidate solutions;
  EVALUATE each candidate;
  REPEAT UNTIL ( TERMINATION CONDITION is satisfied ) DO
    1 SELECT parents;
    2 RECOMBINE pairs of parents;
    3 MUTATE the resulting offspring;
    4 EVALUATE new candidates;
    5 SELECT individuals for the next generation;
  OD
END
```

Fig. 3.1. The general scheme of an evolutionary algorithm in pseudocode

- Components of EAs
 - **representation** (definition of individuals)
 - **evaluation function** (or fitness function)
 - **population**
 - **parent selection mechanism**
 - **variation operators**, recombination and mutation
 - **survivor selection mechanism** (replacement)
 - **termination condition**

- Components of EAs
 - **representation** (definition of individuals)
 - **evaluation function** (or fitness function)
 - **population**
 - **parent selection mechanism**
 - **variation operators**, recombination and mutation
 - **survivor selection mechanism** (replacement)
 - **termination condition**

- **Representation:** What is your genome?
 - **Strings** in a finite alphabet (genetic algorithms)
 - **Real-valued vectors** (evolution strategies)
 - **Finite state machines** (classical evolutionary programming)
 - **Trees** (genetic programming)
- Different approaches may suit different problems better depending on the natural way to encode the candidate solutions

- Components of EAs
 - **representation** (definition of individuals)
 - **evaluation function** (or fitness function)
 - **population**
 - **parent selection mechanism**
 - **variation operators**, recombination and mutation
 - **survivor selection mechanism** (replacement)
 - **termination condition**

- **Evaluation Function**
 - Function that **assigns quality measure to genotypes**
 - **Evaluate phenotype** for quality.
 - Represent requirements that population should adapt to meet (i.e. defines what improvement means)

- Components of EAs
 - **representation** (definition of individuals)
 - **evaluation function** (or fitness function)
 - **population**
 - **parent selection mechanism**
 - **variation operators**, recombination and mutation
 - **survivor selection mechanism** (replacement)
 - **termination condition**

- **Population**
 - holds possible solution
 - A multiset of genotypes (almost always a constant size)
 - Individuals do not change or adapt, population does
 - Selection operators **work at the population level**
 - Best individual **of a given population** is selected

- Components of EAs
 - **representation** (definition of individuals)
 - **evaluation function** (or fitness function)
 - **population**
 - **parent selection mechanism**
 - **variation operators**, recombination and mutation
 - **survivor selection mechanism** (replacement)
 - **termination condition**

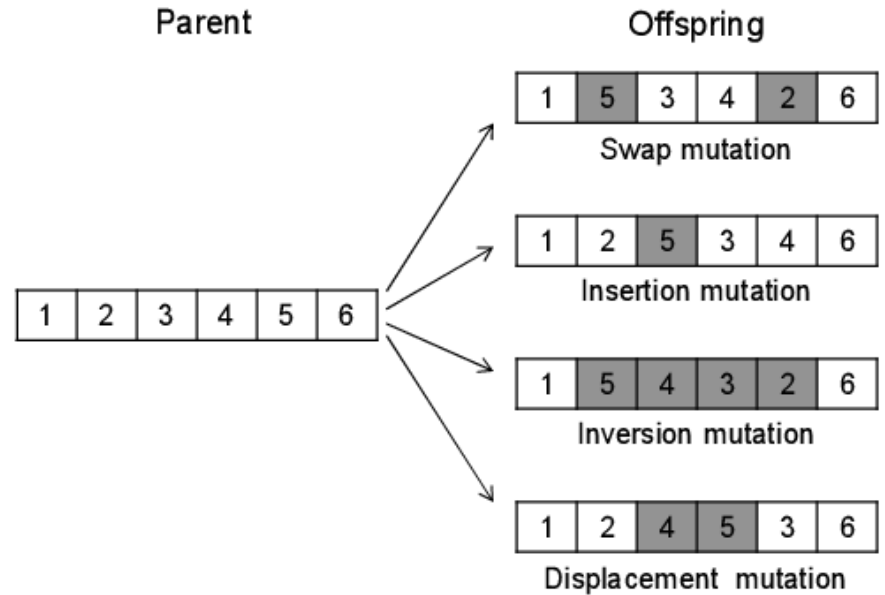
- **Parent selection / mate selection**
 - Distinguish among individuals based on quality
 - Allow better individuals to become parents of next generation
 - Probabalistic – fitter individuals more likely to be selected as parents
 - Often still give low-fitness individuals a small chance, otherwise search becomes too greedy and gets stuck in **local optima**.

- Components of EAs
 - **representation** (definition of individuals)
 - **evaluation function** (or fitness function)
 - **population**
 - **parent selection mechanism**
 - **variation operators, recombination and mutation**
 - **survivor selection mechanism** (replacement)
 - **termination condition**

- **Variation operators**
 - Create new individuals from old ones
 - The “generate” in “generate and test”
 - **Mutation** and **Crossover** (recombination)

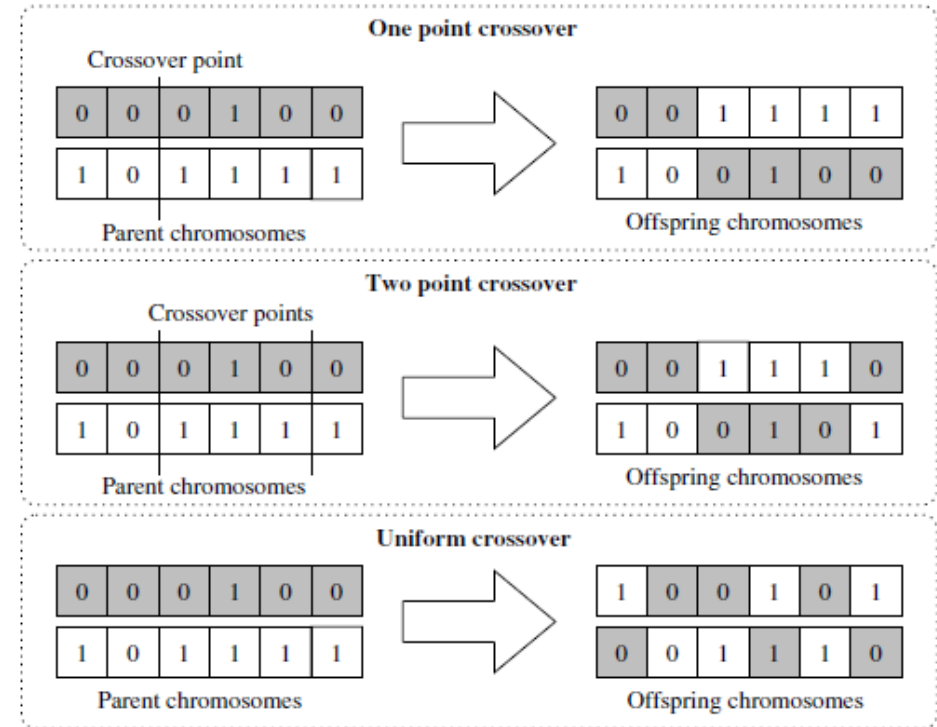
- **Mutation**

- Randomly change part of the genotype
- Stochastic
- Unbiased
- Provides the gene-pool with “fresh blood”



- **Crossover**

- Merges information in two or more parents
- Random recombinations
- Hopefully take good qualities of each parent



- Components of EAs
 - **representation** (definition of individuals)
 - **evaluation function** (or fitness function)
 - **population**
 - **parent selection mechanism**
 - **variation operators**, recombination and mutation
 - **survivor selection mechanism** (replacement)
 - **termination condition**

- **Survivor selection** / replacement strategy
 - Decide which of older generations to keep
 - Based on **fitness values**
 - e.g. keep top % of new and old generation combined
 - Or based on **age**
 - e.g. only keep newest generation

- Components of EAs
 - **representation** (definition of individuals)
 - **evaluation function** (or fitness function)
 - **population**
 - **parent selection mechanism**
 - **variation operators**, recombination and mutation
 - **survivor selection mechanism** (replacement)
 - **termination condition**

- **Termination condition**
 - Stop when we find **optimum fitness** (if you gurantee reaching it)
 - Otherwise
 - 1) Limit **CPU time**
 - 2) Limit total number of **fitness evaluations**
 - 3) **Check if fitness improvement remains low** for a given period of time (i.e. for a number of generations or fitness evaluations)
 - 4) **Check if the population diversity drops** below a given threshold

How do EAs behave?





Fig. 3.4. Typical progress of an EA illustrated in terms of population distribution. For each point x in the search space y shows the corresponding fitness value.

- **Exploration:** generation of individuals in untested space
- **Exploitation:** concentrating in vicinity of known good solutions
- **Premature convergence:** losing population diversity too quickly and getting stuck in a local optimum.

- EAs are **anytime**
 - You can interrupt them at... any time
 - Get the best solution so far, even if suboptimal



- **Improvement is often rapid at start**
- You can initialise your population with **heuristics**, but it might not be worth the effort




Fig. 3.6. Illustration of why heuristic initialisation might not be worth additional effort. Level *a* shows the best fitness in a randomly initialised population; level *b* belongs to heuristic initialisation

- **Improvement slows down**, so long runs may be not worth i t




Fig. 3.7. Why long runs might not be worth performing. *X* shows the fitness increase in the first half of the run, while *Y* belongs to the second half

- Designing and tuning EAs is an art:
 - Population size
 - Number of generations
 - Fitness function
 - Representation
 - Mutation rate
 - Crossover operations
 - Selection procedure
 - Number of solutions to keep


MARIO
003600

0 × 03

WORLD
1-1


TIME
288

Mario Map Generation




Steve Dahlskog and Julian Togelius:
Patterns as Objectives for Level
Generation: PCG Workshop 2013

- Representation attempt 1
 - 2x2 grid of integers storing value for each grid cell
 - Direct mapping to phenotype
 - But “good” space is tiny proportion of search space



- Representation attempt 2
 - A range of integers, each representing the height of an obstacle
 - 0 = hole in the ground
 - 10 = maximum height of obstacle
 - But most obstacles in SMB are close to the ground.

- Representation attempt 3
 - Represent mario level as string of symbols, each symbol corresponds to a column of blocks that existed in the original SMB



- Fitness function
 - Look for level design patterns in the output that match patterns in the original SMB

Table 2: Patterns supported in the fitness function.

Enemies	
Enemy	Low
2-Horde	Low
3-Horde	Low
4-Horde	Low
Roof	Medium
Gaps	
Gaps	Low
Multiple gaps	By stacking
Variable gaps	By stacking
Gap enemy	Low-Medium by stacking
Pillar gap	Pillar High
Valleys	
Valley	Low
Pipe valley	Medium
Empty valley	By stacking
Enemy valley	By stacking
Roof valley	By stacking
Multiple paths	
2-Path	Medium-High
3-Path	Medium-High
Risk and Reward	By stacking
Stairs	
Stair up	Low
Stair down	Low
Empty stair valley	Low
Enemy stair valley	By stacking
Gap stair valley	By stacking

Mario map generation

- Population of 200 (randomly initialised)
- Rank by fitness, top 50% kept
- One point crossover between pairs in rank order (best breeds with second best)
- Mutation: inject a random character in a random position
- 10,000 generations





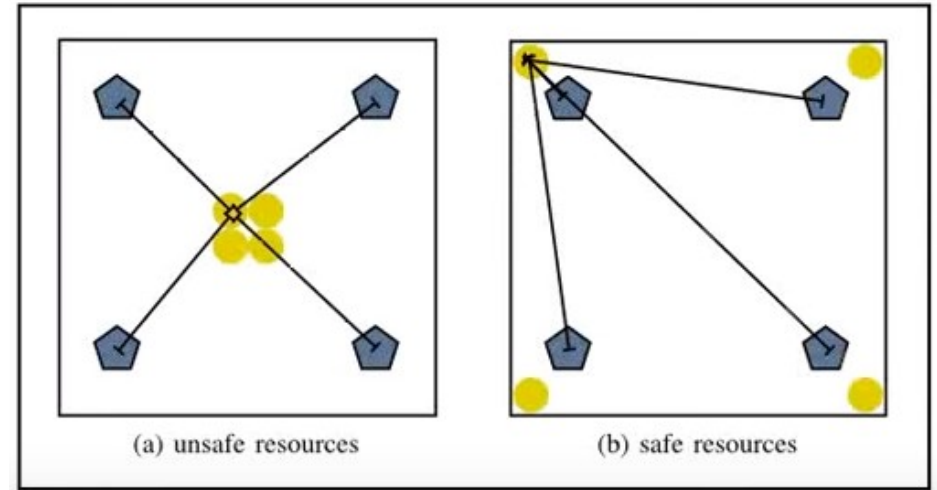
Starcraft Map Generation

Julian Togelius, Mike Preuss, Nicola Beume, Simon Wessing, Johan Hagelbäck, Georgios N. Yannakakis and Corrado Grappiolo (2013):
Controllable Procedural Map Generation via Multiobjective Evolution.
Genetic Programming and Evolvable Machines. Springer.


- Evolution of RTS maps
 - Desirable
 - Playability
 - Fairness
 - Skill differentiation
 - Interestingness



- Fairness fitness function
 - Distance from base to closest resource
 - Resource ownership
 - Resource safety
 - Resource fairness




- Multiobjective EAs have multiple objectives
 - **Pareto front** between them
 - The point at which increasing one objective will decrease another
 - Not always desirable – sometimes you want to look for solutions that are good for **any** of the objectives



PCG iteration as evolution

- 1) Generate some content (**phenotype**) from parameters (**genotype**)
- 2) Score it using designer/AI (**fitness**)
- 3) **Evolve a new set of parameters**
- 4) Repeat



- Eiben & Smith, 2015.
Introduction to Evolutionary Computing (Second Edition)
Springer

