



# Games AI

## Lecture 7.1

Whole Game Search

Muzg the Astute																		20/20	?
1	83	82	81	101	100	102	109	108	143	142	32	31	35	67	66	12	12	38	38


- Whole-Game Search
  - Solving Problems with Serach
  - Tree Search
  - Adversarial Search







- AI Player
  - Turn based games
  - Playing to win



- Turn-based Games
  - Game states
  - Actions



## Introduction





PÉS PARA CIMA, AMANTES  
DAS DIVERSÕES DO WALLY!  
NOSSA! EU PERDI TODAS  
AS MINHAS COISAS, UMA  
EM CADA LUGAR. AGORA  
VOCÊS PRECISAM VOLTAR  
E PROCURÁ-LAS. E, EM  
ALGUM LUGAR, UM DOS  
OBSERVADORES DO WALLY  
PERDEU O POMPOM DE SEU  
CHAPÉU. VOCÊS CONSEGUEM  
VER QUEM É E ENCONTRAR  
O POMPOM QUE FALTA?

Wally



PARA:  
AMANTES DAS  
DIVERSÕES DO WALLY  
DE VOLTA AO INÍCIO,  
COMECE DE NOVO,  
EXCELENTE

## Solving Problems With Search





- Well-defined problems have five components
  - Initial state
  - State-action mapping
  - Transition model
  - Goal test
  - Path cost function


- 1 player game
  - A farmer has a **fox**, a **chicken**, and a **sack of corn**
  - He needs to get them **all across a river**
  - If he leaves the sack of corn with the chicken, **it eats it**
  - If he leaves the fox with the chicken, **it eats it**
  - Only **one thing** can fit in his boat





- **Initial State**

- LeftBank: { Farmer, Chicken, Corn, Fox }
- RightBank: { }




- **State-Action Mapping**
  - What actions can we do from this state?





- **Transition Model**

- What state do we get if we apply action X in state Y?




- **Goal Test**
  - Have we won?



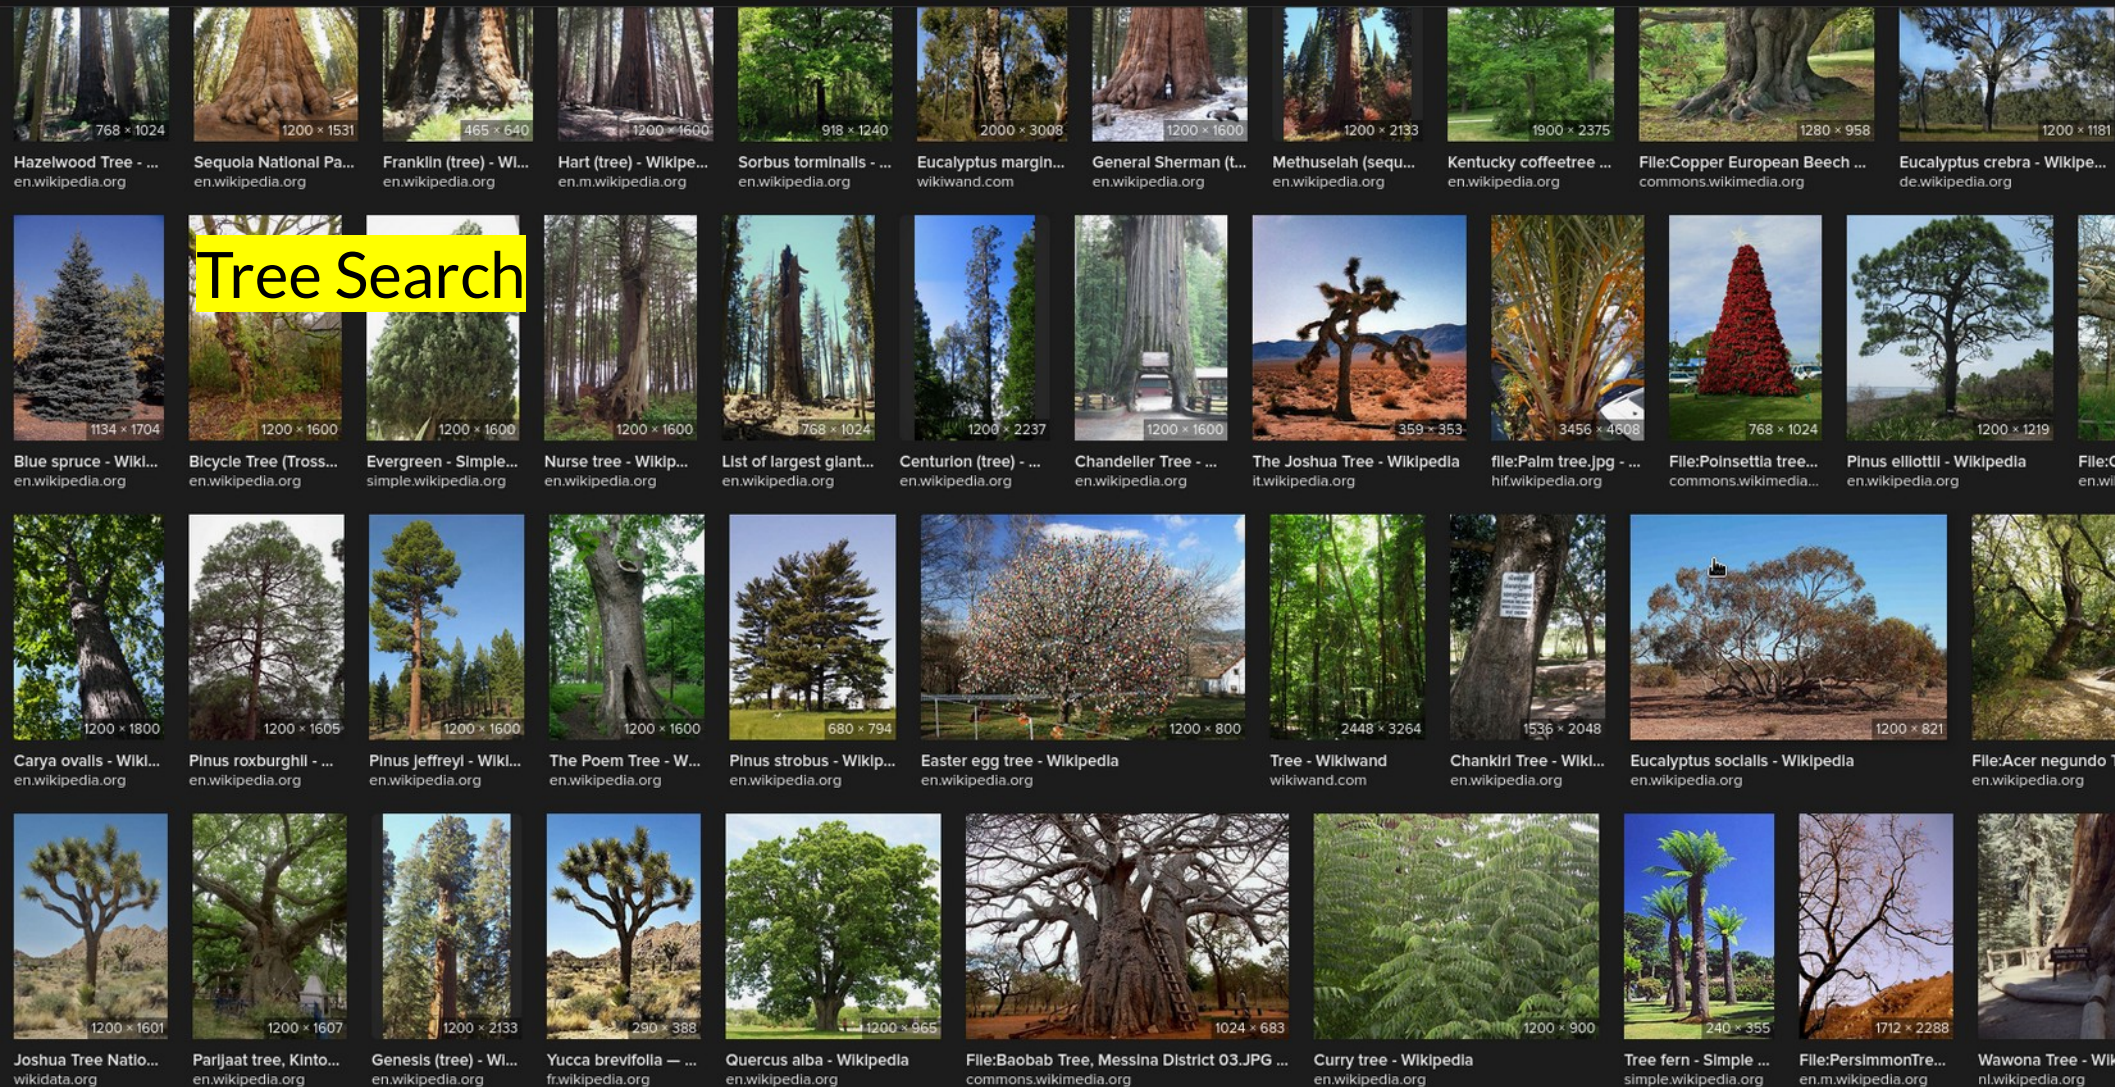


- **Step Cost Function**
  - How expensive is a particular move?




- For many **games**, we can define these five components
  - Initial state
  - State-action mapping
  - Transition model
  - Goal test
  - Path cost function
- Combined, these let us build a **tree of game states**
  - Then we **search** this tree






- We can construct a tree of Game States
  - Root of tree is **Initial State**
  - Each child is the result of **applying one action** (transition model)
  - **Goal states** are leaf nodes
- Let's remind ourselves of a basic tree search algorithm



Queue  
0

- Breadth-first search
- Set up:
  1. Create queue
  2. Add start to queue
  3. Mark as discovered




Queue

- While queue is not empty
  1. Select top of queue
  2. Check if goal




- While queue is not empty
  1. Select top of queue
  2. Check if goal
  3. Add all undiscovered children to queue
  4. Mark each child as discovered




Queue  
1  
2  
3

- While queue is not empty
  1. Select top of queue
  2. Check if goal
  3. Add all undiscovered children to queue
  4. Mark each child as discovered




Queue  
2  
3  
4  
5  
6

- While queue is not empty
  1. Select top of queue
  2. Check if goal
  3. Add all undiscovered children to queue
  4. Mark each child as discovered



Queue  
3  
4  
5  
6  
7  
8  
9

- Depth-first Search
  - Start at the root node
  - Pick the first child
  - Keep exploring down the tree until terminal node is reached.
  - Then backtrack a step and take next branch...





- Depth-first Search Pseudocode


```
procedure DFS_iterative(G, v) is  
  let S be a stack  
  S.push(v)  
  while S is not empty do  
    v = S.pop()  
    if v is not labeled as discovered then  
      label v as discovered  
      for all edges from v to w in G.adjacentEdges(v) do  
        S.push(w)
```

- Depth-first Search
  - Path found may not be shortest
  - Might find a path quickly (depending on tree)
  - Might not terminate (if infinite tree)

- Find a **path to goal state**
  - Path is a sequence of actions = **a plan**
- **Deterministic?**
  - Execute plan step by step
- **Stochastic/Multiplayer?**
  - Execute 1st action in plan
  - Replan

- Guided Search
  - We can use **heuristics** to guide our search (remember  $A^*$ )
  - Heuristics often involve **domain-specific knowledge**
    - e.g. count value of peices in chess
    - What would be a good heuristic for the farmer puzzle?
  - **MCTS** uses **rollouts** (playing the game randomly until win/lose) to decide how interesting a state is
    - It's an example of a **general game-playing** algorithm




A 3D-rendered chessboard with a dark red and black checkered pattern. The board is populated with various chess pieces. Gold pieces are positioned on the left side of the board, while black pieces are on the right. The pieces include pawns, knights, bishops, and a king. The lighting is dramatic, with strong highlights and shadows, giving the pieces a metallic, reflective appearance. The perspective is from a slightly elevated angle, looking down at the board.


Adversarial Search

- Two-player **zero-sum games** (e.g. chess/tic-tac-toe)
  - Two players called **MIN** and **MAX**
- Imagine we're playing MAX
  - On our turn we (MAX) want to pick a **good move**
  - On their turn, MIN wants to pick a **bad move** (for us)
    - Assume MIN is **playing optimally**

- Terminal States
  - If MAX wins, utility = 1
  - If MIN wins, utility = 0




- MINIMAX value of a state  $s$ 
  - MAX's turn: **max** utility of all children
  - MIN's turn: **min** utility of all children






- MINIMAX value of a state  $s$ 
  - MAX's turn: **max** utility of all children
  - MIN's turn: **min** utility of all children




- Which move should we pick?




- The Minimax Algorithm

```
function minimax(node, maximizingPlayer)
  if node is a terminal node then
    return the heuristic value of node
  if maximizingPlayer then
    value :=  $-\infty$ 
    for each child of node do
      value := max(value, minimax(child, FALSE))
    return value
  else /* minimizing player */
    value :=  $+\infty$ 
    for each child of node do
      value := min(value, minimax(child, TRUE))
    return value
```



- Alpha-beta Pruning
  - Sometimes parts of a minimax tree will have no effect on the final outcome
    - We can “prune” those branches (not explore them)


- Say you've got a good move, Move A
  - You check other moves just in case one is better
  - Move B looks pretty good at first, but then you find your opponent can force a win
  - Even if every other outcome of Move B is amazing, you'll never want to take Move B.





# Abstracting for Search

- **Branching Factor**
  - The more moves available, the more search required



- **Depth**
  - The more moves it takes to a goal state, the more search required
    - Use **depth-limited** search, and value non-terminal states using a heuristic

- **Continuous games** still have states - they're just very quick (leading to very deep trees)
  - MCTS can play simple continuous games
  - **Macro actions** = e.g. "Move Left for 10 turns"

- Too complex?
  - Devise strong heuristics
  - Search a **simplified version** of the game
    - Just like we can simplify our pathfinding graph
      - Non-optimal, but faster
  - Cheat

- Summary
  - Construct tree of game states
  - Search tree for path to goal
  - Can do adversarial search
  - Games are often too complex for simple search algorithms
    - Use heuristics
    - Alpha-beta pruning
    - MCTS

